

Task-pressure gradient conflict through training

One seed; 24 bins (29-30 boundaries): median/IQR above, negative-dot fraction below

Naive L1 OL1


$\lambda = 0.05$

$\lambda = 0.1$


$\lambda = 0.5$

$\lambda = 1$

Raw gradient cosine



Conflict boundaries (%)



Training input tokens (billions)