


Pythia-14M: trained endpoints and post-hoc clipping on controls

Uniform TEAL-style clipping is applied only at evaluation; displayed validation loss is capped at 6

- A1-H naive L1
- A4 (trained threshold)
- Post-hoc clipping on ReLU control
- A1-H OL1
- Post-hoc clipping on GeLU control



Post-hoc points above loss 6 are omitted by the explicit display cap (GeLU $p \geq 0.5$; ReLU $p \geq 0.6$).
Lines connect dose/target order only. R_{model} is exact-zero logical-product opportunity, not measured speedup; one seed.