


Pythia-14M: trained A1/A4 endpoints with post-hoc TEAL on every checkpoint

Post-hoc magnitude clipping includes both controls and is evaluation-only; displayed validation loss is capped at 6

- A1-H naive L1 trained endpoints
- A4-Z trained threshold endpoints
- TEAL-augmented nondominated envelope
- A1-H OL1 trained endpoints
- Post-hoc TEAL trajectory (includes GeLU/ReLU controls)



Thin trajectories contain the complete $p=0.0, \dots, 0.9$ sweep; 54 points above loss 6 are omitted only from this view.
Lines connect trained dose or TEAL target order only. R_{model} is exact-zero logical-product opportunity, not measured speedup; one seed.