


Pythia-14M: trained frontiers and post-hoc clipping on controls

A4-OL1 adds four-site orthogonal L1 pressure ($\lambda=1$); displayed validation loss is capped at 6

- A1-H naive L1
- A4 (trained threshold)
- Post-hoc clipping on GeLU control
- A1-H OL1
- A4-OL1 (trained)
- Post-hoc clipping on ReLU control



Post-hoc points above loss 6 are omitted by the explicit display cap (GeLU $p \geq 0.5$; ReLU $p \geq 0.6$).

Lines connect dose/target order only. R_{model} is exact-zero logical-product opportunity, not measured speedup; one seed.