


Pythia-14M: trained frontiers and post-hoc clipping controls

Corrected A4-OL1 uses Run 015; A7 adds post-RoPE Q/K/V gates; A7-OL1 adds seven-site pressure



Post-hoc points above loss 6 are omitted by the inherited source filter (GeLU $p \geq 0.5$; ReLU $p \geq 0.6$); all trained endpoints are shown.
 Lines connect dose/target order only. R_{model} is exact-zero logical-product opportunity, not measured speedup; one seed.