


Logical opportunity vs. measured speedup

Same H100 NVL | BF16 | Uncached 2,048-token prefill

Pythia-70M
Pythia-410M

(a) Batch size 1



A4/A7 denote OL1. Points: paired medians; whiskers: 10th-90th percentiles (7 timing blocks).

Sparse linears only; attention and LM head stay dense. 14M has no valid sparse timing (correctness gate failed).

(b) Batch size 32

