


# The $R_{\text{model}}$ -speed association is model- and implementation-specific

Frozen kernels on one RTX PRO 4500 Blackwell; regression uses only numerically qualified deployments



Each point is one trained checkpoint (one seed). Error bars are 95% input-cluster bootstrap intervals over 16 fixed blocks; 80 paired timings per point.  $R_{\text{model}}$  is a logical-product opportunity, not a runtime metric. OLS lines are descriptive.