



The R_{model} -speed association depends on model shape and GPU



Each point is one trained checkpoint (one seed); bars span independent process estimates. RTX uses 3 eager-only processes; H100 uses 2.
Lines are qualified within-size OLS fits. R_{model} is a logical-product opportunity, not measured runtime savings.