


The same sparse policy changes value across GPU architectures



Lines join the same checkpoint, frozen implementation, batch-one T=2048 BF16 workload, and canonical logical counts.
Speedup is relative to each GPU's own native eager baseline; bars span fresh-process estimates, not training seeds.