

Fresh-process kernel optimization at fixed checkpoint and canonical R_{model}

○ Baseline ● Optimized — Same-repeat pair

Pythia-14M
P0 → K001

Pythia-70M
K009 → K016

Pythia-410M
K004 → K010

Full-model speedup vs. native (x)

A4, $\kappa=.5$
 $R_{model}=12.7\%$

A7, $\kappa=.5$
 $R_{model}=27.5\%$

A4, $\kappa=.5$
 $R_{model}=35.6\%$

A7, $\kappa=.5$
 $R_{model}=40.6\%$

A4, $\kappa=.5$
 $R_{model}=71.6\%$

A7, $\kappa=.5$
 $R_{model}=80.6\%$

Thin links: three counterbalanced process pairs. Large markers: medians. All processes use complete validation.