


Autoresearch progress was architecture-specific and non-monotone



Common score: geometric mean of A4-OL1 $\kappa = 0.5$ and A7-OL1 $\kappa = 0.5$ batch-one full-model speedups on RTX PRO 4500. Vertical spans show the two endpoints. Only milestones with both timings are connected; sample counts vary from 48 to 160 per endpoint. This is one adaptive search trajectory, not optimizer benchmarking.