

(a) Local sparsity at $\kappa = 0.5$ (14M)

FFN (h, m)

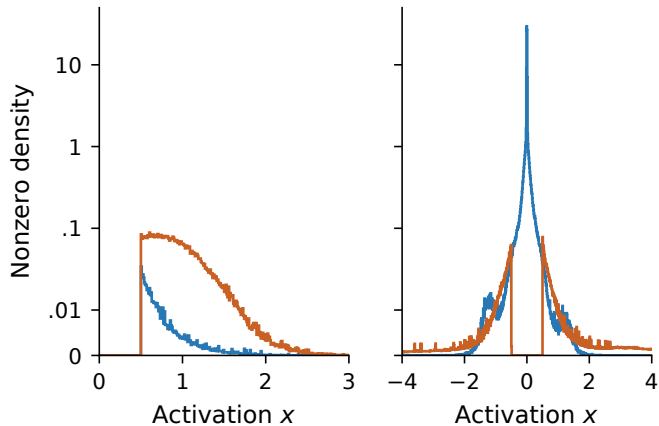
Attention (q, k, v)

4-site + OL1: 99.31% exact zeros

4-site + OL1: 0.22% exact zeros

7-site + OL1: 93.64% exact zeros

7-site + OL1: 95.60% exact zeros



(b) Operation-weighted sparsity accounting

More FFN zeros does not imply higher $\mathcal{S}_{\text{model}}$

QK + PV: 0.06 $\rightarrow$ 16.77 pp

