


(a) Model-wide sparsity predicts realized speedup



(b) Matched kernel ablations

