



Sparsity vs. Quality trade-off and Latency on Pythia 14M and 70M


(a) 14M: quality-sparsity trade-off




(b) 14M: full-model latency



(c) 70M: quality-sparsity trade-off



(d) 70M: full-model latency



○ Base model
● GeLU → ReLU

---●--- T_4/P_h ---●--- T_7/P_h
---●--- T_4/P_{all} ---●--- T_7/P_{all}