



# Quality, sparsity and full-model latency on Pythia-14M

40 trained models | Latency: RTX 5090, BF16, B=1, T=2048 | three timing sessions

(a) Quality-sparsity trade-off



(b) Full-model latency



- $\bigcirc$   $T_0/P_0$  (Base)
- $\bullet$   $T_1/P_0$  (ReLU)
- $\text{---}\bullet\text{---}$   $T_1/P_1$
- $\text{---}\bullet\text{---}$   $T_{hz}/P_h$  (new)

- $\text{---}\bullet\text{---}$   $T_4/P_0$
- $\text{---}\bullet\text{---}$   $T_4/P_h$
- $\text{---}\bullet\text{---}$   $T_4/P_{\text{all}}$
- $\text{---}\bullet\text{---}$   $T_7/P_0$
- $\text{---}\bullet\text{---}$   $T_7/P_h$
- $\text{---}\bullet\text{---}$   $T_7/P_{\text{all}}$