

Quality-latency trade-offs across model scales

