

Targeted sparsification exhibits scale-dependent quality-latency trade-offs

