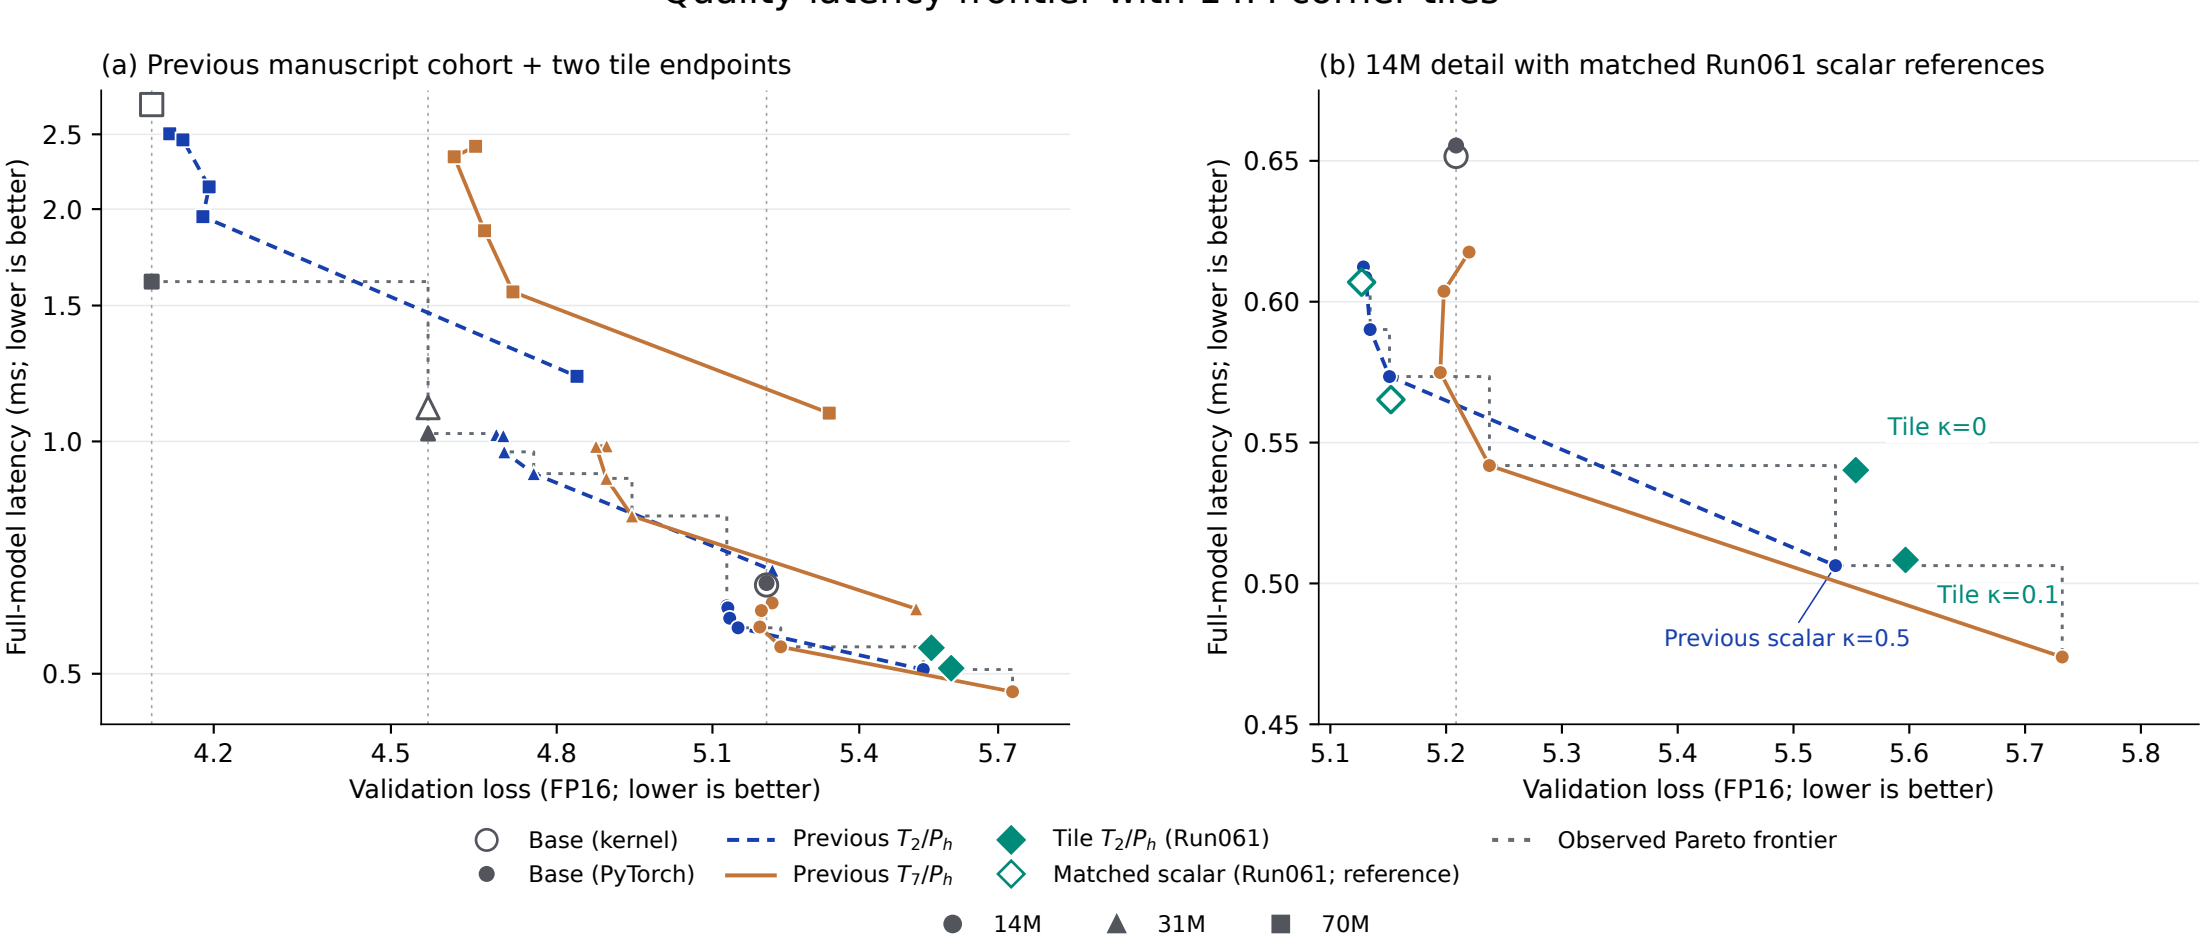
Quality-latency frontier with 14M corner tiles
(a) Previous manuscript cohort + two tile endpoints
(b) 14M detail with matched Run061 scalar references
Full-model latency (ms; lower is better)
Validation loss (FP16; lower is better)
2.5
2.0
1.5
1.0
0.5
4.2
4.5
4.8
5.1
5.4
5.7
0.65
0.60
0.55
0.50
0.45
5.1
5.2
5.3
5.4
5.5
5.6
5.7
5.8
Tile κ=0
Tile κ=0.1
Previous scalar κ=0.5
Base (kernel)
Base (PyTorch)
Previous T_2/P_h
Previous T_7/P_h
Tile T_2/P_h (Run061)
Matched scalar (Run061; reference)
Observed Pareto frontier
14M
31M
70M
Frontier: 36 previous points + 2 tiles; hollow scalar references excluded.
Whiskers span 3 process means. Sessions/kernels differ; raw coordinates, no normalization.