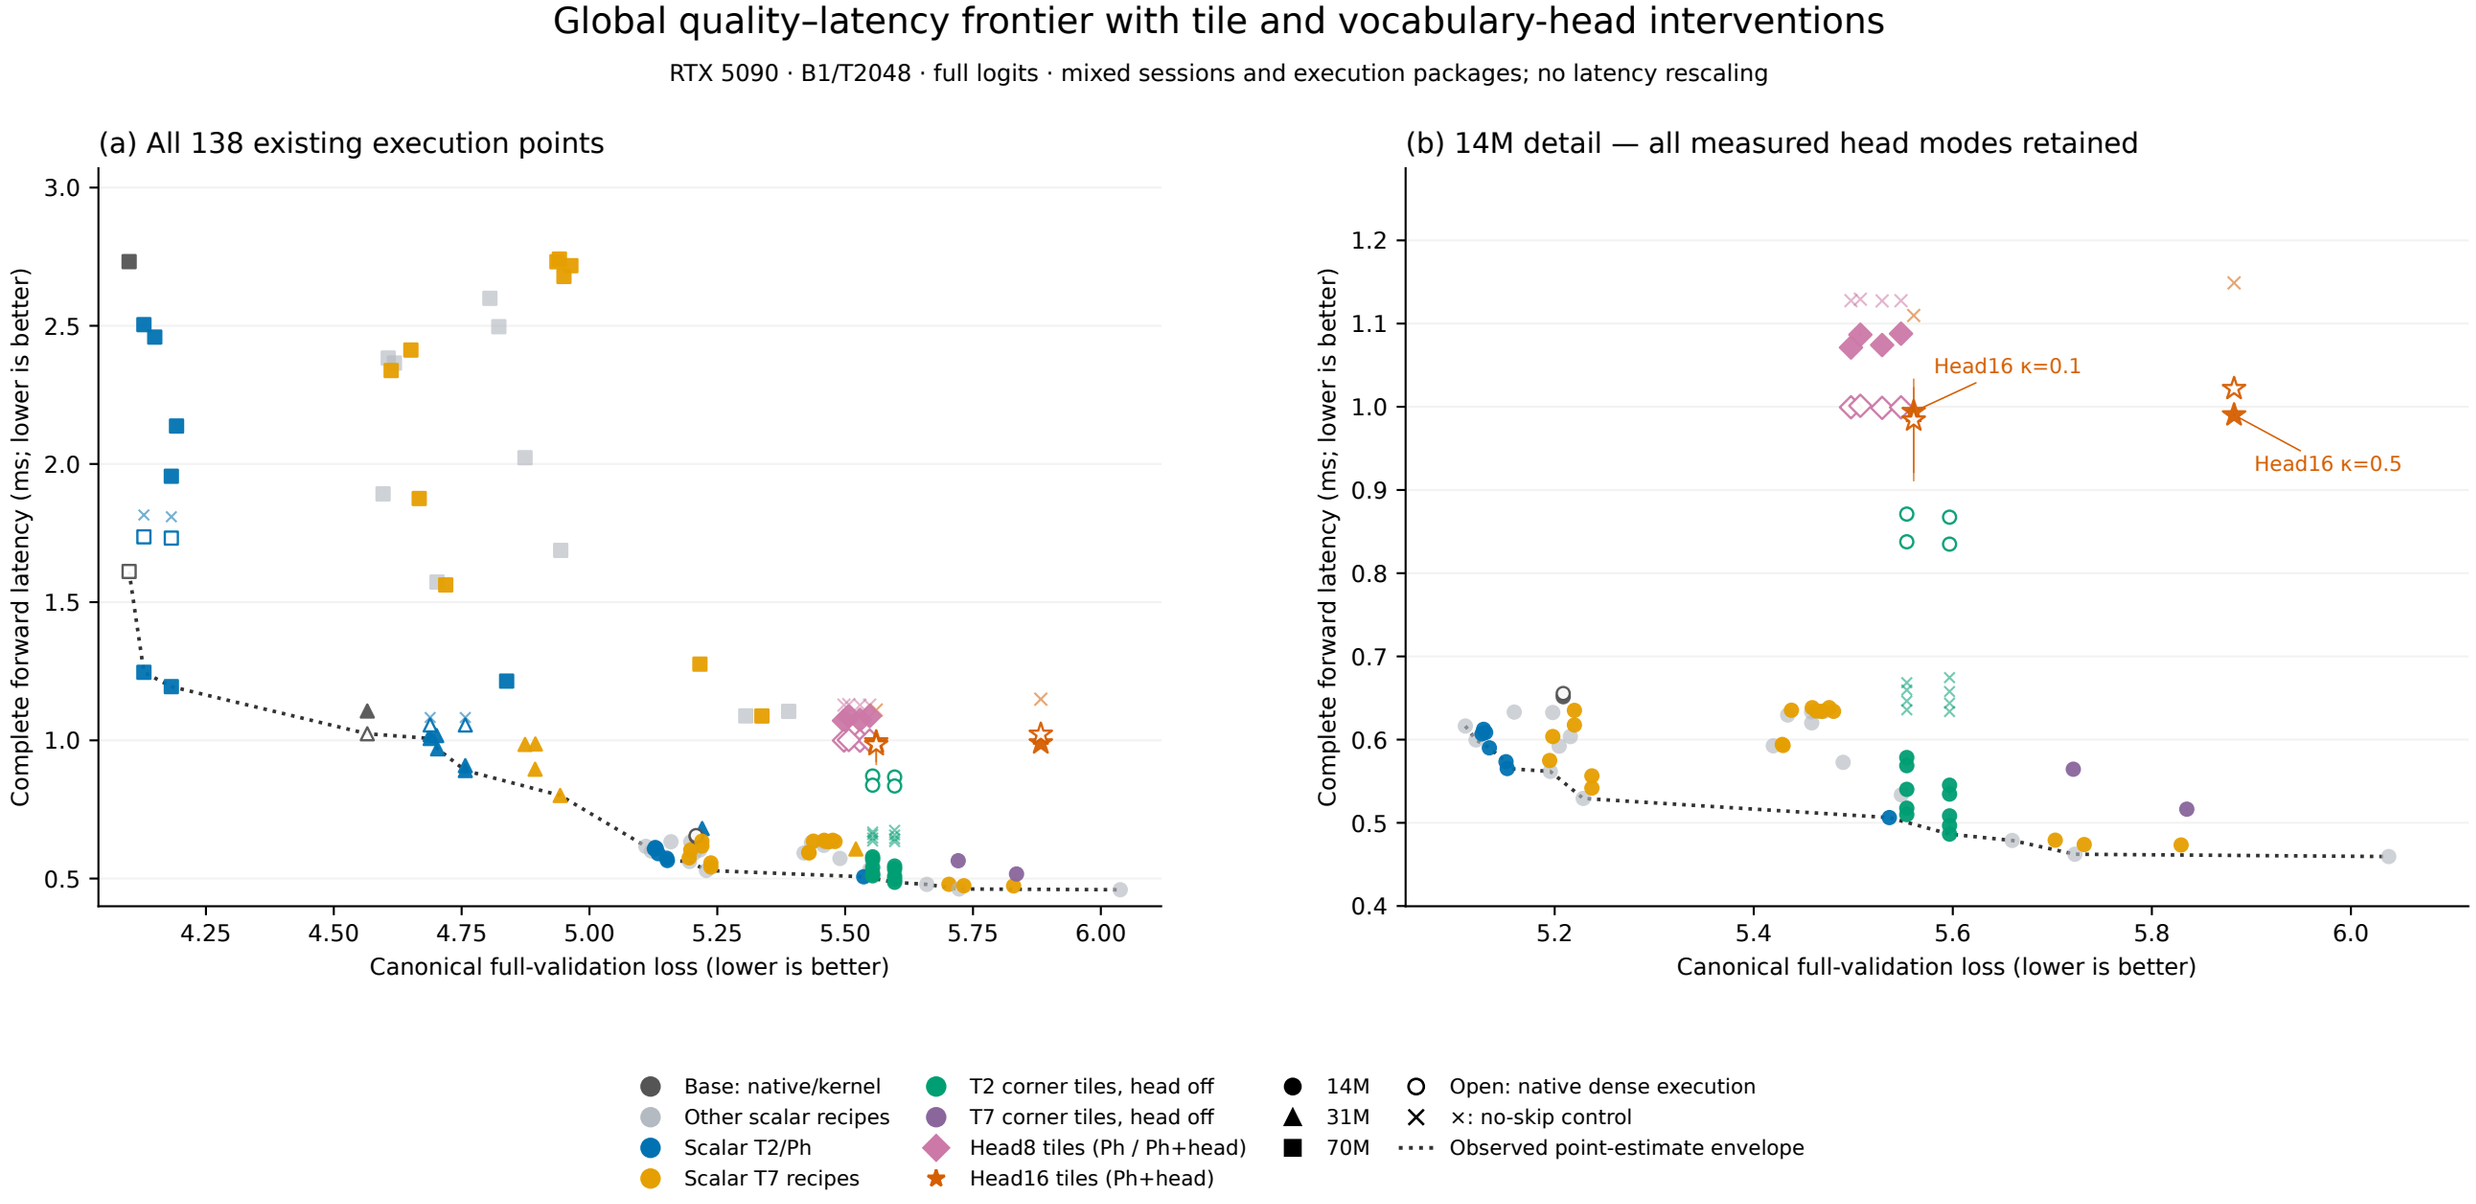
Global quality–latency frontier with tile and vocabulary-head interventions
RTX 5090 · B1/T2048 · full logits · mixed sessions and execution packages; no latency rescaling
(a) All 138 existing execution points
(b) 14M detail — all measured head modes retained
Complete forward latency (ms; lower is better)
Canonical full-validation loss (lower is better)
Head16 κ=0.1
Head16 κ=0.5
Base: native/kernel
Other scalar recipes
Scalar T2/Ph
Scalar T7 recipes
T2 corner tiles, head off
T7 corner tiles, head off
Head8 tiles (Ph / Ph+head)
Head16 tiles (Ph+head)
14M
31M
70M
Open: native dense execution
×: no-skip control
Observed point-estimate envelope