

ACTIVATION SPARSITY AS TRAINING–INFERENCE CO-DESIGN

Anonymous authors

Paper under double-blind review

ABSTRACT

Activation sparsity can accelerate transformer inference when a specialized kernel can skip computation associated with zero activations. We ask how broadly and aggressively train-time sparsification should be applied to obtain such exploitable sparsity without excessive quality loss. Using matched Pythia-family pretraining runs, we systematically vary the placement and intensity of activation pressure and fixed-threshold nonlinearities, and connect the resulting sparsity to specialized-kernel execution. We find that broader interventions can substantially increase model-wide sparsity without reducing latency: at 14M parameters, more than doubling model-wide sparsity from 12.71% to 27.48% slightly increases latency. At 14M, conditional kernel ablations identify the FFN hidden activation h and attention-output projection input z as the most useful sparsification sites. Pressure applied only at h also increases sparsity at z , despite z receiving no direct pressure, while broader pressure produces stronger optimization conflict at less responsive sites. These observations motivate thresholding only h, z while applying pressure only at h . At 14M, this intervention achieves lower validation loss than the dense control while reducing latency from 0.655 to 0.573 ms. Comparisons at 31M and 70M retain the lower loss of targeted versus seven-site thresholding at matched thresholds, but not the validation-loss improvement over the dense control. Our results suggest that sparsification is best treated as a training–inference co-design problem: rather than maximizing sparsity globally, training should induce specific exploitable structure for the inference system.

1 INTRODUCTION

Activation sparsity can improve transformer inference efficiency when specialized kernels can skip computation and memory accesses of zero activations. Recently, Cetin et al. (2026) has successfully demonstrated inference speedups from train-time pressure and specialized kernel design. This result motivates a follow-up question: *how broadly and aggressively should train-time interventions be applied throughout the transformer to induce exploitable sparsity while reducing latency at an acceptable quality cost?*

Existing approaches differ in where and how they induce sparsity. TEAL applies post-hoc, training-free thresholding to pretrained models; ProSparse and Cetin et al. (2026) apply ℓ_1 regularization to feed-forward (FFN) activations during training; and Q-Sparse applies top- k sparsification to both FFN and attention activations (Liu et al., 2024; Song et al., 2024; Wang et al., 2024). However, these studies do not isolate the effects of different train-time sparsity interventions under a common training protocol. Under matched pretraining conditions using models from the Pythia family (Biderman et al., 2023), we systematically vary the coverage and intensity of train-time sparsity interventions across activation sites and measure their model-wide sparsity (S_{model}), and execution efficiency.

We define an intervention by the placement of pressure (P) and thresholding (T). To keep the intervention space tractable, for P we introduce Adam-orthogonalized ℓ_1 activation pressure (OL1), limiting conflict with the task update while avoiding an additional pressure-strength sweep. For T we use a fixed-threshold nonlinearity with $\kappa \in \{0, 0.01, 0.05, 0.1, 0.5\}$. The P and T placements are independent, so a recipe such as T_7/P_{all} specifies their respective site sets, while sweeping κ controls sparsification aggressiveness and traces the recipe’s trade-off curve. Each model is then connected with an specialized kernel that connects its sparsity directly with execution benefit.




Figure 1: **Sparsity, execution savings and quality on Pythia-14M.** (a) Model-wide logical sparsity S_{model} versus full-model latency. The arrow connects T_7/P_{all} to T_7/P_h at $\kappa = 0.5$: 27.48% and 16.66% sparsity yield similar latency. (b) Conditional execution ablations on T_7/P_{all} at $\kappa = 0.5$ identify h, z as the profitable skipping sites. Saved time is latency with one skipping path disabled minus latency with all paths enabled. (c) Validation loss versus latency exposes the quality cost of broader sparsification. Panels (a,c) show the same 40 checkpoints: T_2/P_h is blue, T_7/P_{all} orange, and other recipes gray. Hollow markers and vertical guides denote base dense model; lines connect settings within recipes. Timings use BF16 on an RTX 5090, batch one and 2,048 tokens.

Under extensive 14M ablations, we first observe that broader intervention coverage can substantially extend attainable model-wide sparsity without translating into lower latency. At $\kappa = 0.5$, expanding pressure and thresholding from four to seven sites more than doubles model-wide sparsity, from 12.71% to 27.48%, while latency slightly increases from 0.459 to 0.473 ms (Fig. 1(a)). Conditional kernel ablations then identify h and z as the profitable latency skipping paths, with the largest benefit at h (Fig. 1(b)). Sparsity measurements reveal a related nonlocal effect: applying pressure only at h also increases sparsity at the thresholded z site, despite z receiving no direct pressure (Fig. 3). We call this effect *sparsity spillover*. OL1 adds an additional geometry evidence for h : applying pressure beyond h increases both the pressure-update norm and the component removed by orthogonalization, with the strongest effect at q, k, v , indicating that these sites are less susceptible to sparsification (Fig. 4).

Together, these ablations point to h and z as the most useful sparsification sites. They are more responsive to intervention and their sparsity translates into latency reduction, even when model-wide sparsity is smaller than broader recipes. This motivates a targeted T_2/P_h intervention: apply pressure only at h , exploit its spillover to z , and threshold only h and z . At 14M and $\kappa \leq 0.1$, T_2/P_h achieves both lower validation loss than the dense control and lower full-model latency (0.573 vs. 0.655 ms, Fig. 1(c)). We further compare T_2/P_h and T_7/P_h at 31M and 70M to test whether narrowing threshold coverage remains favorable while holding pressure placement fixed at h . The combinatorial space of train-time interventions and kernel designs motivates inverting the design problem: first specify the structure the kernel can exploit, then design train-time interventions to induce it.

2 RELATED WORK

2.1 SPARSIFICATION INTERVENTIONS

Post-hoc thresholding removes small activations at inference-time of a pre-trained model. While CATS thresholds small-magnitude activations within FFNs (Lee et al., 2024), TEAL extends it to attention (Liu et al., 2024). These studies motivate training-time interventions to adapt activation distributions to sparse computation and improve the trade-off between quality and sparsity.

ReLU Strikes Back studies the effect of activations functions in language models (Mirzadeh et al., 2023), showing clear sparsity gains on the use of ReLUs. ProSparse combines progressive ℓ_1 pressure and thresholding at FFNs during continued training (Song et al., 2024). It does not, however, extend its interventions to attention and restrict the analysis to local FFN sparsity. Cetin et al. (2026) extends ProSparse setup to pre-training and propose specialized kernels. These studies show that training can reshape the quality–sparsity trade-off, but they’re restricted to a single local intervention at FFN sites.

Q-Sparse extend sparsification beyond FFNs. It applies top- k selection to both attention and FFN projection inputs (Wang et al., 2024), but with no ℓ_1 pressure. Its top- k operation introduces, however, additional overhead due to the sorting operation. Spark Transformer tries to address it with statistical top- k (You et al., 2025), but we take a simpler approach and use a constant thresholding non-linearity. We argue that when the intervention is introduced during training, the network can adapt itself. Additionally, we evaluate the complementary effect of regularization pressure on both attention and FFNs.

2.2 INFERENCE SPEEDUP

Turning activation sparsity into speedup requires a specialized kernel adapted to the model architecture. TEAL accelerates batch-one autoregressive decoding, reporting up to $1.8\times$ speedup at 50% sparsity (Liu et al., 2024). Cetin et al. (2026) address batched matrix multiplication with their TwELL kernel, and report $1.3\times$ speedup for batch inference. Spark likewise couples its sparsity mechanism with an implementation that reports decoding gains (You et al., 2025). In each setting, skipped computation and memory access must compensate the additional overhead introduced by the sparse operations. This development has been increasingly accelerated by the use of coding agents, through iterative kernel development with execution and profiling feedback (Ouyang et al., 2025; Luo et al., 2026). We take a related approach and use GPT-6 Astra to develop a specialized kernel based on our training interventions (OpenAI, 2026).

3 METHODOLOGY

3.1 SPARSIFICATION INTERVENTIONS

A sparsification intervention specifies where activations are thresholded (thresholding sites $\mathcal{N}$) and where activation pressure is applied (pressure sites $\mathcal{P}$). The sets $\mathcal{N}$ and $\mathcal{P}$ need not coincide, and the thresholding and pressure functions may take different forms. Together, these choices define a combinatorial space of training interventions. We restrict this space to keep the computational cost manageable while addressing the central research questions.

For thresholding, we avoid the overhead of exact top- k selection (Wang et al., 2024; You et al., 2025) and instead use one-sided and symmetric fixed-threshold operators $G_{+,\kappa}(x) = x\mathbf{1}\{x \geq \kappa\}$ and $G_{\pm,\kappa}(x) = x\mathbf{1}\{|x| \geq \kappa\}$, with $\kappa \geq 0$. The constant κ controls the strength of the intervention and is held fixed within each training condition. At $\kappa = 0$, $G_{+,\kappa}$ reduces to ReLU, whereas $G_{\pm,\kappa}$ reduces to the identity. We use one-sided thresholding at projection inputs, where signed weight products can recover either sign, and symmetric thresholding for attention operands to better preserve their signed dot-product distribution.

For activation pressure, ℓ_1 averages across targeted sites and layers the mean absolute activation at each site after any thresholding. Ordinary L1 pressure minimizes $\mathcal{L}_{\text{task}} + \lambda\ell_1$, where λ controls the pressure intensity (Song et al., 2024; Cetin et al., 2026). To limit conflict between the ℓ_1 pressure and the task update while reducing the need to tune an additional hyperparameter, we introduce orthogonal ℓ_1 pressure (OL1), inspired by gradient-conflict projection methods (Yu et al., 2020; Hsieh et al., 2024). OL1 adds a separately preconditioned pressure correction after the task-only AdamW update, projects away its component opposing the adaptive task direction, and caps its norm relative to that direction. At a fixed optimizer state, once the norm budget is active, further increasing λ leaves the pressure correction unchanged up to numerical stabilizers. To keep the multisite intervention study tractable, we use a common $\lambda = 1$ and budget $b = 1$ and directly measure whether the budget becomes active during training. Appendix A.1 gives the full formulation, a brief comparison and shows that budget saturation depends strongly on the target set.

Figure 2 locates the intervention sites within the Pythia architecture. The variables a and m denote the normalized inputs to the attention and FFN branches, respectively; h is the FFN hidden activation; z is the concatenated attention context before the output projection; and q , k , and v are the query, key, and value operands. Attention probabilities are denoted by p . We apply $G_{+,\kappa}$ at selected sites in $\{a, m, h, z\}$ and $G_{\pm,\kappa}$ at selected sites in $\{q, k, v\}$. When h is targeted, the thresholding nonlinearity replaces GELU. Table 1 summarizes the studied thresholding and pressure assignments and their sparsity ceilings across model sizes. Section 3.2 defines these ceilings.

Table 1: **Intervention recipes and analytic sparsity ceilings.** T_1 replaces GeLU with ReLU at h . Pressure uses OL1. For $T_2/T_4/T_7$, we sweep $\kappa \in \{0, .01, .05, .1, .5\}$ with $\lambda = 1$ and $b = 1$; the matched 70M latency sweep for T_2/P_h stops at $\kappa = 0.1$, with its additional $\kappa = 0.5$ endpoint reported separately in Appendix D.1; the 14M T_1/P_1 ablation instead sweeps $\lambda \in \{.05, .1, .5, 1\}$ at $b = 1$. Ceilings are $100 \cdot \mathcal{S}_{\text{model}}^{\max}$ at 2,048 tokens. These ceilings quantify architectural reach, not whether every recipe was executed at every size. Executed conditions are listed in Appendix C.3.

Recipe	Threshold sites $\mathcal{N}$	Pressure sites $\mathcal{P}$	Ceiling (%)			
			14M	31M	70M	410M
T_0/P_0 (Base Model)	$\emptyset$	$\emptyset$	0.00	0.00	0.00	0.00
T_1/P_0	h	$\emptyset$	4.28	7.58	12.35	24.93
T_1/P_1	h	h	4.28	7.58	12.35	24.93
T_2/P_h	h, z	h	5.35	9.48	15.44	31.16
T_4/P_0	a, m, h, z	$\emptyset$	12.83	22.75	37.06	74.78
T_4/P_h	a, m, h, z	h	12.83	22.75	37.06	74.78
T_4/P_{all}	a, m, h, z	$\mathcal{N}$	12.83	22.75	37.06	74.78
T_7/P_0	a, m, h, z, q, k, v	$\emptyset$	29.95	37.92	49.42	87.25
T_7/P_h	a, m, h, z, q, k, v	h	29.95	37.92	49.42	87.25
T_7/P_{all}	a, m, h, z, q, k, v	$\mathcal{N}$	29.95	37.92	49.42	87.25




Figure 2: **Intervention sites in a shared Pythia block.** The attention and FFN branches use separate LayerNorms and join the residual path. Green nodes denote activation sites, and purple nodes denote learned weights. Table 1 specifies the thresholding and pressure assignments.

3.2 SPARSITY MEASURE AND ARCHITECTURAL REACH

Activation sparsity is the fraction of zero activations at a site. We define *model-wide sparsity* as $\mathcal{S}_{\text{model}} = N_0/N_{\text{model}}$, where N_0 counts scalar multiplications with a zero activation operand in Pythia-block matrix products, and N_{model} counts all scalar multiplications in those blocks and the final output projection. We aggregate counts over validation data, excluding masked attention pairs.

To quantify intervention potential a priori, we define the *model-wide sparsity ceiling* as $\mathcal{S}_{\text{model}}^{\max} = N_{\text{reach}}(\mathcal{N})/N_{\text{model}}$, where $N_{\text{reach}}(\mathcal{N})$ counts scalar multiplications guaranteed a zero activation operand if all selected thresholding sites are 100% sparse. Thus, $\mathcal{S}_{\text{model}}^{\max}$ is the maximum sparsity directly attributable to $\mathcal{N}$. It depends on architecture, selected sites, and sequence length. Appendix B gives the precise accounting and definitions.

4 EXPERIMENTAL STUDY

We pretrain Pythia-family models at 14M, 31M, 70M and 410M from random initialization (Biderman et al., 2023) on MiniPile (Kaddour, 2023). Extensive 14M ablations vary intervention placement; a cross-scale comparison evaluates targeted and seven-site thresholding at 14M, 31M and 70M.

Additional 70M recipes and a 410M fixed-token stress test appear in Appendix D. Each run receives 712 optimizer steps and 1.493 billion input tokens. Within each size, runs share initialization, data order and training budget. We evaluate loss and $\mathcal{S}_{\text{model}}$ on the full validation set. Appendix C gives the protocol.

4.1 EXTENSIVE ABLATIONS AT 14M

We compare interventions at 14M, separating aggressive thresholding ($\kappa = 0.5$) from moderate settings ($\kappa = \{0.05, 0.1\}$). First, **broader coverage increases attainable sparsity without consistently reducing latency**. At $\kappa = 0.5$, T_7/P_{all} more than doubles the sparsity of T_4/P_{all} (27.48% versus 12.71%) at 0.473 versus 0.459 ms (Table 2). Within T_7 , broadening pressure from P_h to P_{all} adds 10.82 percentage points of sparsity with nearly unchanged latency, while loss rises from 5.732 to 5.829. All seven aggressive recipes reduce latency relative to Base but increase loss.

Table 2: **Sparsity and latency under aggressive thresholding**. All seven thresholded 14M recipes at $\kappa = 0.5$. All latencies, including Base, use the specialized 14M kernel on RTX 5090 (BF16, batch one, 2,048 tokens, full logits), with three timing processes per checkpoint.

Recipe	Ceiling (%)	Sparsity (%)	Latency (ms)	Loss
Base (kernel)	0.00	0.00	0.652	5.209
T_2/P_h	5.35	5.34	0.506	5.536
T_4/P_0	12.83	10.22	0.479	5.660
T_4/P_h	12.83	10.23	0.462	5.723
T_4/P_{all}	12.83	12.71	0.459	6.038
T_7/P_0	29.95	15.39	0.479	5.703
T_7/P_h	29.95	16.66	0.474	5.732
T_7/P_{all}	29.95	27.48	0.473	5.829

Moderate thresholding with local pressure offers a better quality-latency balance. At both moderate thresholds, P_h gives lower loss and latency than P_{all} under both T_4 and T_7 (Table 3). At $\kappa = 0.05$, T_4/P_h reaches 0.562 ms with loss 5.196, versus Base’s 0.652 ms and 5.209. Narrowing thresholding to T_2/P_h improves loss further to 5.134.

Table 3: **Quality-latency trade-offs at moderate thresholds**. All seven thresholded 14M recipes at $\kappa \in \{0.05, 0.1\}$, with the same metrics and timing protocol as Table 2.

Recipe	κ	Ceiling (%)	Sparsity (%)	Latency (ms)	Loss
Base (kernel)	—	0.00	0.00	0.652	5.209
T_2/P_h	0.05	5.35	5.00	0.590	5.134
	0.1	5.35	5.17	0.573	5.151
T_4/P_0	0.05	12.83	8.21	0.630	5.434
	0.1	12.83	8.95	0.593	5.420
T_4/P_h	0.05	12.83	9.04	0.562	5.196
	0.1	12.83	9.34	0.529	5.229
T_4/P_{all}	0.05	12.83	10.46	0.573	5.490
	0.1	12.83	11.39	0.534	5.548
T_7/P_0	0.05	29.95	9.13	0.635	5.438
	0.1	29.95	10.43	0.594	5.429
T_7/P_h	0.05	29.95	10.13	0.575	5.195
	0.1	29.95	10.90	0.542	5.237
T_7/P_{all}	0.05	29.95	9.86	0.635	5.463
	0.1	29.95	11.80	0.593	5.429

4.2 SPILLOVER: LOCAL PRESSURE INDUCES NONLOCAL SPARSITY

Pressure at h increases sparsity at z under both T_4 and T_7 . At $\kappa = 0.05$, P_h raises h sparsity from 74–85% to 94–99% across layers (Figure 3). Sparsity at the unpressured z site also rises in every layer: from 71–85% to 83–92% under T_4 , and from 70–87% to 84–94% under T_7 . This spillover is therefore not confined to four-site thresholding, and suggests that z is responsive to the changes induced by pressure at h . This is evidence of the network representation globally re-adapting itself from a local intervention.




Figure 3: **Activation sparsity by site and layer in Pythia-14M.** The six threshold/pressure recipes at $\kappa = 0.05$: T_4 (top) and T_7 (bottom), each under P_0 , P_h and P_{all} . Each cell gives the percentage of zero activations over all complete validation sequences at one site and layer. Values are rounded to whole percentages, so 0 or 100 does not establish a completely dense or all-zero tensor.

Broader pressure reveals unequal sparse susceptibility across sites. Under T_4/P_{all} , extending pressure beyond h increases sparsity across the targeted sites, with z reaching 99.74–99.85% sparsity in layers 2–4. Extending pressure further under T_7/P_{all} produces a different pattern: sparsity at q, k, v increases but remains only 3–13%, while h and z become less sparse than under P_h in every layer. The OL1 geometry in Figure 4 is consistent with this asymmetry. Pressure outside h produces larger pre-cap correction norms and larger components opposing the task update, and the effect is strongest for T_7/P_{all} . At $\kappa = 0.05$, the median removed component rises from 0.15–0.17% under P_h to 1.12% for T_4/P_{all} and 66.84% for T_7/P_{all} , while the corresponding pre-cap norm ratios rise from 0.08–0.10 to 0.28 and 46.00. Broadening pressure increasingly triggers OL1 projection and norm capping, potentially reducing the effective pressure reaching h and z . Together with the spillover effect, these results indicate that h and z are more susceptible to sparsification than the additional sites.

4.3 WHEN SPARSITY TRANSLATES TO SPEEDUP

We next test whether the sites with the most susceptibility to sparsity also respond most strongly to execution time. Table 4 isolates each skipping path at 14M by disabling it while keeping all the other paths fixed. It reports site-sparsity, *bypass* (the percentage of omitted matrix instructions) and conditional saved time for T_7/P_{all} at $\kappa = 0.5$. We test over this recipe because its the sparser variant. The result is aligned with the susceptibility observation: **the measurable savings comes only from h and z .**

Sparsity reduces latency when avoided work exceeds the overhead of detecting and skipping it. At h, z , the sparse kernel checks 8×16 activation tiles (i.e. sub-matrices) before loading their associated weights, so empty tiles avoid both weight reads and matrix multiplication. Applying the same strategy




Figure 4: **OL1 geometry depends on pressure placement at 14M.** (a) Median removed opposing component, $100\rho_{\text{opp}}$; bars span the 25th–75th percentiles over 712 training steps. (b) Pre-cap norm ratio r/b ; bold curves show pointwise median over individual threshold curves. The dotted line marks the norm cap.

Table 4: **Conditional latency savings:** 14M T_7/P_{all} , $\kappa = 0.5$. Latency is measured per 2,048-token sequence on RTX 5090 (BF16, batch one). Saved time is full-model latency with one skipping path off minus latency with all paths on; positive values indicate a benefit (reduced latency). Brackets give [minimum, maximum] savings across all off/on pairs of three independent timing runs. Each run reports the geometric mean over 64 sequences.

Operation	Activation sparsity (%)	Bypass (%)	Saved time (μs)	Process span (μs)
a : QKV projection	75.6	8.4	−1.6	[−6.0, +3.8]
m : FFN-up	68.6	0.0	−2.4	[−5.1, +2.7]
h : FFN-down	99.9	94.5	+150.1	[+145.4, +154.4]
z : attention-output projection	99.9	99.6	+29.0	[+23.7, +35.7]
q, k : QK scores	93.5 / 94.5	58.0	−2.9	[−7.3, +1.1]
v : PV product	98.7	67.2	−6.2	[−9.0, −2.2]

Activation zeros pooled over layers and validation sequences; q/k reported separately.

at a and m doesn’t have an expressive bypass, and instead increased full-model latency by 17.2% and 27.4%, because the skipped work does not surpass the additional overhead. We speculate **that the relation between local sparsity and bypass is not linear, and only over a critical level it translates in a structure that produces high-zero tiles with high probability.**

Attention operations q, k, v follows FlashAttention’s tiled computation (Dao et al., 2022). In our kernel implementation, the zero checks happens only after loading the operands. Skipping a matrix instruction therefore avoids the arithmetic without avoiding those reads and softmax operation. Additionally, softmax turns zero scores into nonzero probabilities, so q, k sparsity does not automatically propagate to p . Enabling qk/pv skipping together raises full-model latency from 0.514 to 0.521 ms despite substantial bypass. So, for this implementation, attention’s arithmetic savings do not offset its overhead, while h/z benefit from avoiding weight reads as well as computation.

4.4 TARGETED AND BROAD THRESHOLDING ACROSS MODEL SIZES

Holding pressure placement fixed across scales. Figure 5 compares T_2/P_h and T_7/P_h at 14M, 31M and 70M. Both apply OL1 pressure only at h , while threshold coverage changes from h, z to all seven sites. The intermediate 31M experiment extends this controlled comparison of threshold placement, rather than replicating the full 14M intervention study.

Targeted thresholding preserves quality, but not uniformly relative to Base. At every matched κ , T_2/P_h has lower validation loss than T_7/P_h at all three sizes. Broader thresholding nevertheless reaches lower minimum latency in each sweep, so neither recipe is uniformly preferable across both metrics. Moderate T_2/P_h improves loss over Base at 14M, but all evaluated 31M and 70M settings incur a positive loss penalty. At 31M, $\kappa = 0.1$ —the same moderate threshold used in the main 14M comparison—gives loss 4.757 and latency 0.908 ms, versus Base loss 4.566 and native latency 1.024 ms.




Figure 5: **Quality-latency trade-offs across model scales.** Circles, triangles and squares denote 14M, 31M and 70M. Dashed blue T_2/P_h thresholds h, z ; solid orange T_7/P_h thresholds all seven sites. Both apply OL1 pressure only at h . Segments connect measured points in increasing $\kappa \in \{0, 0.01, 0.05, 0.1, 0.5\}$ order; they do not represent measured interpolated frontiers. Complete results appear in Appendix D.1. Hollow and filled gray markers show Base under kernel and native PyTorch execution; labeled vertical guides mark Base loss. Both axes are logarithmic. Timings use RTX 5090, BF16, batch one, 2,048 tokens and full vocabulary logits.

Within-recipe latency reductions differ from speedups over native execution. Increasing κ from 0 to 0.1 within T_2/P_h reduces kernel latency by 0.039 ms at 14M, 0.112 ms at 31M and 0.549 ms at 70M. These are changes within each implementation, not speedups over native Base. At 70M, $\kappa = 0.1$ remains slower than native Base; $\kappa = 0.5$ is faster but incurs a larger loss increase (Table 8).

Kernel efficiency remains size-dependent. Base takes 1.106 ms under the specialized implementation at 31M, versus 1.024 ms natively. At 70M, the partially optimized implementation takes 2.732 ms versus 1.611 ms natively. The training-side trade-off must therefore be evaluated with the implementation used at each size. The intermediate-scale result broadens the evidence for targeted thresholding, but does not establish automatic kernel portability or quality-free speedups.

5 DISCUSSION AND CONCLUSION

Our results distinguish the benefits of targeted thresholding from the costs of broader pressure. At 14M, execution ablations and optimization diagnostics identify h, z as useful sparsification sites under the studied implementation. Comparisons of T_2/P_h and T_7/P_h across 14M, 31M and 70M further show that narrowing threshold coverage preserves validation quality at matched thresholds, while broader thresholding can reach lower latency at a larger quality cost. The validation-loss improvement over the dense control is confined to the evaluated 14M conditions. The broader result is a favorable

quality–latency trade-off from targeting useful sites, rather than a general improvement in quality from sparsification.

More broadly, sparsification is one instance of a training–inference co-design problem. The relevant question is not simply where a model can be made sparse, but which representation makes expensive computation removable at inference time. A systems-guided approach is therefore to identify the operations that dominate latency, determine the structured property that would make them cheaper or skippable, and enforce that property during training. This principle may extend beyond activation sparsity to tiled activation or weight sparsity, quantization, batched execution, and other hardware-dependent forms of structure.

5.1 FAILURES AND LIMITATIONS

Our main analysis combines extensive 14M intervention ablations with targeted-versus-broad thresholding comparisons at 14M, 31M and 70M. We also trained 410M models, but the fixed-token budget places that size in a different optimization regime: with the same token budget, the 410M base model has higher loss than 70M and larger gradient norms (Appendix D.2). Kernel efficiency also does not transfer directly across scales. Adapting the 14M kernel to 70M introduces substantial overhead; further optimization improves execution, but does not close the gap with native PyTorch, although it benefits from the sparse structure. Our runtime study is deliberately narrow (batch-one, full-sequence inference on a single GPU), aimed to provide a controlled setting for studying the interaction between train-time sparsification and execution, rather than a complete production-serving evaluation. Training comparisons use one seed per condition. This limits claims about individual point estimates, although the consistent trade-off structure across dozens of matched recipes provides evidence for the broader trends we report. Finally, the intervention space remains only partially explored. The 31M comparison holds pressure placement fixed at h ; it does not independently establish the spillover mechanism, site-wise optimization conflict or conditional execution savings measured at 14M. Pressure and threshold placements are combinatorial, and broad-pressure experiments change several pressure targets at the same time. Because OL1 also averages pressure across targeted sites, these comparisons do not isolate the contribution of each added site. Future work could separate pressure targets one at a time, remove this cross-site averaging confound, and allow site-specific thresholds rather than sharing a single κ across sites.

AI USE STATEMENT

We used OpenAI Codex with GPT-6 Astra to assist with experimental design refinement, mathematical exposition, implementation and debugging of training and evaluation code, GPU-kernel development, experiment execution and monitoring, analysis and interpretation of results, figure preparation, and manuscript drafting and revision. The authors directed the research and approved experimental designs. AI-assisted implementations were checked through tests and comparisons with reference computations; the kernel validation procedures and their limitations are documented in Appendix E. The authors take responsibility for the final text, code, results and scientific claims.

REFERENCES

- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. Pythia: A Suite for Analyzing Large Language Models Across Training and Scaling. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 2397–2430. PMLR, 2023. URL <https://proceedings.mlr.press/v202/biderman23a.html>.
- Edoardo Cetin, Stefano Peluchetti, Emilio Castillo, Akira Naruse, Mana Murakami, and Llion Jones. Sparser, Faster, Lighter Transformer Language Models. *arXiv preprint arXiv:2603.23198*, 2026. URL <https://arxiv.org/abs/2603.23198>.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. *arXiv preprint arXiv:2205.14135*, 2022. URL <https://arxiv.org/abs/2205.14135>.

- Yu-Guan Hsieh, James Thornton, Eugene Ndiaye, Michal Klein, Marco Cuturi, and Pierre Ablin. Careful with that Scalpel: Improving Gradient Surgery with an EMA. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 19085–19100. PMLR, 2024. URL <https://proceedings.mlr.press/v235/hsieh24a.html>.
- Jean Kaddour. The MiniPile Challenge for Data-Efficient Language Models. *arXiv preprint arXiv:2304.08442*, 2023. URL <https://arxiv.org/abs/2304.08442>.
- Donghyun Lee, Je-Yong Lee, Genghan Zhang, Mo Tiwari, and Azalia Mirhoseini. CATS: Contextually-Aware Thresholding for Sparsity in Large Language Models. In *Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=v3w2a7EInO>.
- James Liu, Pragaash Ponnusamy, Tianle Cai, Han Guo, Yoon Kim, and Ben Athiwaratkun. Training-Free Activation Sparsity in Large Language Models. *arXiv preprint arXiv:2408.14690*, 2024. URL <https://arxiv.org/abs/2408.14690>.
- Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. In *International Conference on Learning Representations*, 2019. URL <https://arxiv.org/abs/1711.05101>.
- Mao Luo, Hongbin Li, Feng Lin, Hanling Yi, and Zhe Huang. Agentic Kernel Optimization: Generating State-of-the-Art GPU Kernels Without Hand-Written CUDA. *arXiv preprint arXiv:2608.14560*, 2026. URL <https://arxiv.org/abs/2608.14560>.
- Iman Mirzadeh, Keivan Alizadeh, Sachin Mehta, Carlo C. Del Mundo, Oncel Tuzel, Golnoosh Samei, Mohammad Rastegari, and Mehrdad Farajtabar. ReLU Strikes Back: Exploiting Activation Sparsity in Large Language Models. *arXiv preprint arXiv:2310.04564*, 2023. URL <https://arxiv.org/abs/2310.04564>.
- OpenAI. GPT-6 Astra: A New Generation of Intelligence. OpenAI, September 2026. URL <https://openai.com/index/gpt-6-astra/>. Accessed September 19, 2026.
- Anne Ouyang, Simon Guo, Simran Arora, Alex L. Zhang, William Hu, Christopher Ré, and Azalia Mirhoseini. KernelBench: Can LLMs Write Efficient GPU Kernels? In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 47356–47415. PMLR, 2025. URL <https://proceedings.mlr.press/v267/ouyang25a.html>.
- Chenyang Song, Xu Han, Zhengyan Zhang, Shengding Hu, Xiyu Shi, Kuai Li, Chen Chen, Zhiyuan Liu, Guangli Li, Tao Yang, and Maosong Sun. ProSparse: Introducing and Enhancing Intrinsic Activation Sparsity within Large Language Models. *arXiv preprint arXiv:2402.13516*, 2024. URL <https://arxiv.org/abs/2402.13516>.
- Hongyu Wang, Shuming Ma, Ruiping Wang, and Furu Wei. Q-Sparse: All Large Language Models can be Fully Sparsely-Activated. *arXiv preprint arXiv:2407.10969*, 2024. URL <https://arxiv.org/abs/2407.10969>.
- Chong You, Kan Wu, Zhipeng Jia, Lin Chen, Srinadh Bhojanapalli, Jiaxian Guo, Utku Evci, Jan Wassenberg, Praneeth Netrapalli, Jeremiah J. Willcock, Suvinay Subramanian, Felix Chern, Alek Andreev, Shreya Pathak, Felix Yu, Prateek Jain, David E. Culler, Henry M. Levy, and Sanjiv Kumar. Spark Transformer: Reactivating Sparsity in FFN and Attention. *arXiv preprint arXiv:2506.06644*, 2025. URL <https://arxiv.org/abs/2506.06644>.
- Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. Gradient Surgery for Multi-Task Learning. In *Advances in Neural Information Processing Systems*, 2020. URL <https://arxiv.org/abs/2001.06782>.

A INTERVENTION DEFINITIONS

A.1 ACTIVATION PRESSURE AND OL1

Objective and update. For J activation tensors A_j at the sites and layers targeted by $\mathcal{P}$, each with n_j entries after any thresholding, define

$$\ell_1(\theta; \mathcal{P}) = \frac{1}{J} \sum_{j=1}^J \frac{1}{n_j} \sum_{i=1}^{n_j} |A_{j,i}|. \quad (1)$$

Each tensor has equal weight, so changing $\mathcal{P}$ changes both the targets and normalization. Ordinary L1 applies AdamW (Loshchilov & Hutter, 2019) to the globally clipped gradient of $\mathcal{L}_{\text{task}} + \lambda \ell_1$. OL1 separates task and pressure updates: task-only AdamW is followed by a pressure correction that removes the opposing component and limits its norm. Pressure therefore does not directly enter Adam’s moment estimates.

Both gradients are evaluated at the same parameters and averaged over the same microbatches. The clipped task gradient updates AdamW’s bias-corrected moments $\hat{m}_t, \hat{v}_t$; the pressure gradient $g_1 = \nabla_{\theta} \ell_1$ is unclipped and has no momentum. Using the same preconditioner for both directions gives

$$u = \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon_A}, \quad w = \frac{g_1}{\sqrt{\hat{v}_t} + \epsilon_A}. \quad (2)$$

Norms and inner products include parameters with both gradients and initialized Adam state. For $c = \langle u, w \rangle$, $q = \|u\|_2^2$ and stabilizer $\epsilon > 0$,

$$\tilde{w} = \begin{cases} w - \frac{c}{q + \epsilon} u, & c < 0 \text{ and } q > \epsilon, \\ w, & \text{otherwise.} \end{cases} \quad (3)$$

Aligned pressure is retained. For weight $\lambda \geq 0$ and relative norm budget $b > 0$,

$$r = \frac{\lambda \|\tilde{w}\|_2}{\|u\|_2 + \epsilon}, \quad s = \min \left\{ 1, \frac{b}{r + \epsilon} \right\} \quad (r > 0), \quad (4)$$

with $s = 1$ when $r = 0$. After task-only AdamW, each parameter group p receives $\theta_p \leftarrow \theta_p - \alpha_p \lambda s \tilde{w}_p$, where α_p is its learning rate. Before learning-rate scaling and ignoring stabilizers, the correction norm is $\min\{\lambda \|\tilde{w}\|_2, b \|u\|_2\}$. At a fixed optimizer state, increasing λ has no further effect once the cap is active.

OL1 combines conflict projection (Yu et al., 2020) and a smoothed task direction (Hsieh et al., 2024) in Adam’s preconditioned coordinates. The stabilizer makes orthogonality approximate. This calculation excludes group learning rates and decoupled weight decay; avoiding opposition to the Adam direction does not guarantee unchanged task or validation loss.

L1 versus OL1. At 14M T_1/P_1 , matched initialization, data order and training budget yield similar results (Table 5). Sparsity differs by less than 0.02 percentage points at every λ . OL1 has slightly lower loss at $\lambda \in \{0.05, 0.1, 0.5\}$ and higher loss at $\lambda = 1$. One seed cannot establish whether these small loss differences persist across seeds.

Table 5: **L1 and OL1 at T_1/P_1 on Pythia-14M.** Final validation loss and model-wide sparsity after 712 matched training steps; OL1 uses $b = 1$. Both metrics cover all 338 complete validation blocks (500 documents; 1,444 tail tokens excluded). One initialization per condition.

λ	L1		OL1	
	S_{model} (%)	Val. loss	S_{model} (%)	Val. loss
0.05	3.1416	5.2061	3.1452	5.1981
0.1	3.3363	5.1655	3.3386	5.1594
0.5	3.7876	5.1127	3.7680	5.1102
1	3.9493	5.1023	3.9384	5.1210

B SPARSITY ACCOUNTING

B.1 ACTIVATION AND MODEL-WIDE SPARSITY

Activation sparsity is the fraction of zero entries at a site, pooling counts across layers and validation sequences before dividing.

Model-wide sparsity counts scalar multiplications with a zero activation input:

$$\mathcal{S}_{\text{model}} = \frac{N_0}{N_{\text{model}}}. \quad (5)$$

N_0 counts these multiplications within Pythia blocks; N_{model} counts all multiplications in those blocks and the final vocabulary projection. The latter contributes only to the denominator. Counts are pooled across layers and validation sequences; percentages are $100\mathcal{S}_{\text{model}}$.

For a block’s linear projection from d_{in} to d_{out} features over n rows, the denominator contribution is $nd_{\text{in}}d_{\text{out}}$. Each zero input contributes d_{out} to N_0 , one multiplication per output feature. Zero weights do not count as activation sparsity.

In attention, both operands are activations. We count each multiplication once if either or both are zero. For queries q , keys k , values v and attention probabilities p , one head in one sequence contributes

$$N_{0,\text{QK}} = \sum_{t=1}^T \sum_{u=1}^t \sum_{j=1}^{d_h} \mathbf{1}\{q_{tj} = 0 \vee k_{uj} = 0\}, \quad (6)$$

$$N_{0,\text{PV}} = \sum_{t=1}^T \sum_{u=1}^t \sum_{j=1}^{d_h} \mathbf{1}\{p_{tu} = 0 \vee v_{uj} = 0\}. \quad (7)$$

Here T is sequence length, d_h the head dimension, t the query position and u the key/value position. Restricting $u \leq t$ excludes masked future positions from both counts. We sum over heads, layers and validation sequences, using the actual operands after thresholding and intervening transformations.

Normalization, softmax, activation functions, additions and memory transfers are excluded, so $\mathcal{S}_{\text{model}}$ does not directly measure latency or speedup.

B.2 MODEL-WIDE SPARSITY CEILING

Assume all activations at the thresholding sites $\mathcal{N}$ are zero. If $N_{\text{reach}}(\mathcal{N})$ counts the multiplications that must then have a zero activation input, the ceiling is

$$\mathcal{S}_{\text{model}}^{\text{max}} = \frac{N_{\text{reach}}(\mathcal{N})}{N_{\text{model}}}. \quad (8)$$

The denominator is unchanged, and each multiplication is counted once. The numerator includes downstream activations forced to zero: setting $v = 0$ also makes $z = 0$, affecting both PV and the attention-output projection. Zero attention scores do not imply zero softmax probabilities.

The ceiling depends only on architecture, selected sites and sequence length. Natural zeros outside those sites can raise measured $\mathcal{S}_{\text{model}}$ above it; the ceiling does not predict the associated quality cost. Appendix C.2 gives the counts for each Pythia size.

C EXPERIMENTAL SETUP DETAILS

C.1 PYTHIA ARCHITECTURE AND INTERVENTION SITES

Table 6 gives the shapes and positions of the sites in Figure 2. Here T is sequence length, d hidden width, d_f FFN width and H the number of heads, each with dimension $d_h = d/H$. Sites and threshold are shared across layers.

Table 6: Activation shapes and intervention positions for one sequence.

Site	Shape	Following matrix product	Position
a	$T \times d$	QKV projection	After attention LayerNorm
m	$T \times d$	FFN up projection	After FFN LayerNorm
h	$T \times d_f$	FFN down projection	Activation-function output
q, k	$H \times T \times d_h$	Attention scores	After RoPE
v	$H \times T \times d_h$	Probability-value	After the QKV split
z	$T \times d$	Attention-output projection	Concatenated attention context

C.2 ARCHITECTURE-SPECIFIC SPARSITY CEILINGS

For one sequence, each Pythia block uses $3Td^2$ scalar multiplications in the QKV projection, Td^2 in the attention-output projection, Tdd_f in each FFN projection, and $dT(T+1)/2$ in each of QK and PV . With L layers and vocabulary size V_{vocab} , the totals are

$$\begin{aligned} C_{\text{block}} &= T(4d^2 + 2dd_f) + dT(T+1), \\ C_{\text{model}} &= LC_{\text{block}} + TdV_{\text{vocab}}. \end{aligned} \quad (9)$$

Sequence count cancels from the ceiling ratio. For $d_f = 4d$, per-layer counts in the ceiling numerator are 0 for T_0/P_0 , $4Td^2$ for T_1/P_0 and T_1/P_1 , $12Td^2$ for T_4 , and C_{block} for T_7 . For T_2 , the per-layer numerator is $Tdd_f + Td^2$, covering the FFN-down and attention-output projections; with $d_f = 4d$, this becomes $5Td^2$. Multiplying by L/C_{model} gives $\mathcal{S}_{\text{model}}^{\text{max}}$, independent of pressure placement.

Table 1 uses $V_{\text{vocab}} = 50,304$ and $(L, d) = (6, 128), (6, 256), (6, 512), (24, 1,024)$ for 14M, 31M, 70M and 410M, respectively, with the common sequence length in Table 7.¹

The 31M configuration has six layers, hidden width 256, FFN width 1,024, eight heads of dimension 32, vocabulary size 50,304 and 30,494,720 trainable parameters.

Table 7: **Training settings across model sizes.** Effective batch size equals microbatch size times gradient accumulation. The common training budget gives fewer tokens per parameter at larger sizes.

Setting	14M	31M	70M	410M
Parameters	14,067,712	30,494,720	70,426,624	405,334,016
Peak learning rate	10^{-3}	10^{-3}	10^{-3}	3×10^{-4}
Minimum learning rate	10^{-4}	10^{-4}	10^{-4}	3×10^{-5}
Microbatch (sequences)	32	4	4	4
Gradient accumulation (microbatches)	32	256	256	256
Effective batch (sequences/update)	1,024	1,024	1,024	1,024
Sequence length	2,048	2,048	2,048	2,048
Tokens per parameter	106.14	48.96	21.20	3.68

C.3 TRAINING AND EVALUATION

All sizes use the Pythia-14M-deduped tokenizer and seeds of 1234 for initialization and data order, with one seed per condition. AdamW settings are shared: $\beta = (0.9, 0.95)$, $\epsilon = 10^{-8}$, weight decay 0.1 and gradient clipping at norm 1. Attention and hidden dropout are zero. Training uses 712 updates with 1% linear warmup followed by cosine learning-rate decay. Table 7 gives the learning rates, microbatch sizes and shared effective batch. Gradients accumulate over the listed microbatches before each optimizer update.

The 31M experiment comprises Base and T_2/P_h , T_7/P_h at each $\kappa \in \{0, 0.01, 0.05, 0.1, 0.5\}$. Both thresholded recipes use OL1 at h with $\lambda = b = 1$. All eleven runs share the same random initial weights, training-block order, optimizer, learning-rate schedule, token budget and validation coverage. Base and T_2/P_h come from Run054; T_7/P_h comes from Run055, with identical initialization and data-order hashes. Training uses FP16 with dynamic loss scaling and FP32 parameters and optimizer

¹Pinned architecture configurations: 14M, 31M, 70M, and 410M.




Figure 6: **Base-model (T_0/P_0) training at 14M, 31M, 70M and 410M.** (a) Training loss and (b) gradient norm before clipping (logarithmic scale). Faint lines show individual updates; solid lines show centered nine-update means, with shorter windows at the ends. Gradient norms are computed after accumulation and are not normalized by parameter count.

state; AdamW excludes biases and normalization parameters from weight decay. These runs extend the threshold-placement comparison with pressure fixed at h .

At 14M, the executed conditions include Base, T_1/P_0 , the T_1/P_1 pressure-weight ablation and all seven thresholded recipes in Table 1. At 70M, they include Base, T_1/P_0 , T_2/P_h , and T_4, T_7 with P_h or P_{all} ; the additional T_2/P_h , $\kappa = 0.5$ endpoint was measured in a separate session. The 410M conditions are Base, T_1/P_0 and the five-threshold T_4/P_{all} , T_7/P_{all} sweeps. Validation uses all 500 MiniPile documents: 338 complete 2,048-token blocks (692,224 input tokens), excluding the final 1,444-token tail. Published losses use the canonical training evaluation; the BF16 kernel checks below are a separate numerical qualification.

Figure 6 shows the training trajectories. They do not establish convergence, and the lower learning rates at 410M limit comparisons that attribute loss differences to model size alone.

D COMPLETE RESULTS AND SUPPORTING DIAGNOSTICS

D.1 DETAILED TRAINED RESULTS

Table 8 compares the recipes evaluated at both 14M and 70M, including settings outside the favorable quality–latency trade-off. Loss and model-wide sparsity use the same definitions and final checkpoints as the main figures, with one training seed and complete validation per condition. Loss and latency deltas are relative to each size’s dense PyTorch Base. Table 9 adds the 31M Base and matched T_2/P_h , T_7/P_h sweeps.

Table 8: **14M/70M results.** Parentheses give changes in loss and latency relative to each size’s T_0/P_0 Base. Base uses dense PyTorch/SDPA; other rows use the optimized implementation for their size. Deltas use the fixed Base references shown, with negative values indicating reductions. Latency is in ms (RTX 5090, BF16, batch one, 2,048 tokens, full logits). OL1 uses $\lambda = b = 1$.

Recipe	κ	14M			70M		
		$\mathcal{S}_{model}$ (%)	Loss (Δ)	Latency (Δ)	$\mathcal{S}_{model}$ (%)	Loss (Δ)	Latency (Δ)
T_0/P_0 (Base)	—	0.00	5.21 (+0.000)	0.66 (+0.000)	0.00	4.10 (+0.000)	1.61 (+0.000)
T_1/P_0 (ReLU)	—	2.71	5.27 (+0.061)	0.63 (-0.023)	10.07	4.22 (+0.123)	2.72 (+1.106)

Recipe	κ	14M			70M		
		$\mathcal{S}_{\text{model}}$ (%)	Loss (Δ)	Latency (Δ)	$\mathcal{S}_{\text{model}}$ (%)	Loss (Δ)	Latency (Δ)
T_2/P_h	0	4.63	5.13 (-0.080)	0.61 (-0.043)	14.63	4.13 (+0.029)	2.50 (+0.893)
	0.01	4.72	5.13 (-0.078)	0.61 (-0.047)	14.70	4.15 (+0.050)	2.46 (+0.848)
	0.05	5.00	5.13 (-0.074)	0.59 (-0.065)	15.03	4.19 (+0.093)	2.14 (+0.526)
	0.1	5.17	5.15 (-0.057)	0.57 (-0.082)	15.21	4.18 (+0.083)	1.96 (+0.345)
	0.5	5.34	5.54 (+0.328)	0.51 (-0.149)	15.43	4.84 (+0.738)	1.21 (-0.397)
T_4/P_h	0	8.41	5.22 (+0.007)	0.60 (-0.052)	25.79	4.61 (+0.507)	2.38 (+0.772)
	0.01	8.59	5.20 (-0.004)	0.59 (-0.063)	25.97	4.62 (+0.519)	2.37 (+0.754)
	0.05	9.04	5.20 (-0.013)	0.56 (-0.094)	26.57	4.60 (+0.497)	1.89 (+0.281)
	0.1	9.34	5.23 (+0.020)	0.53 (-0.126)	27.22	4.70 (+0.602)	1.57 (-0.038)
	0.5	10.23	5.72 (+0.514)	0.46 (-0.193)	29.24	5.31 (+1.206)	1.09 (-0.523)
T_4/P_{all}	0	7.81	5.46 (+0.250)	0.63 (-0.022)	25.67	4.81 (+0.706)	2.60 (+0.988)
	0.01	8.59	5.46 (+0.250)	0.62 (-0.035)	26.22	4.82 (+0.723)	2.50 (+0.886)
	0.05	10.46	5.49 (+0.281)	0.57 (-0.083)	28.49	4.87 (+0.774)	2.02 (+0.412)
	0.1	11.39	5.55 (+0.340)	0.53 (-0.122)	30.57	4.94 (+0.844)	1.69 (+0.076)
	0.5	12.71	6.04 (+0.829)	0.46 (-0.196)	35.60	5.39 (+1.290)	1.10 (-0.506)
T_7/P_h	0	8.40	5.22 (+0.011)	0.62 (-0.038)	25.79	4.65 (+0.551)	2.41 (+0.801)
	0.01	8.86	5.20 (-0.010)	0.60 (-0.052)	26.01	4.61 (+0.512)	2.34 (+0.727)
	0.05	10.13	5.19 (-0.014)	0.57 (-0.081)	26.89	4.67 (+0.567)	1.88 (+0.264)
	0.1	10.90	5.24 (+0.029)	0.54 (-0.114)	27.78	4.72 (+0.619)	1.56 (-0.049)
	0.5	16.66	5.73 (+0.523)	0.47 (-0.182)	32.23	5.34 (+1.237)	1.09 (-0.523)
T_7/P_{all}	0	7.05	5.48 (+0.272)	0.63 (-0.021)	23.56	4.94 (+0.841)	2.74 (+1.130)
	0.01	7.72	5.48 (+0.267)	0.64 (-0.017)	24.23	4.94 (+0.837)	2.73 (+1.121)
	0.05	9.86	5.46 (+0.254)	0.63 (-0.021)	26.52	4.96 (+0.864)	2.72 (+1.106)
	0.1	11.80	5.43 (+0.221)	0.59 (-0.062)	28.69	4.95 (+0.850)	2.68 (+1.067)
	0.5	27.48	5.83 (+0.621)	0.47 (-0.182)	40.60	5.22 (+1.116)	1.28 (-0.335)

The broader-pressure comparison remains distinct: at $\kappa = 0.5$, T_7/P_{all} raises loss above Base by 1.116 at 70M and 0.621 at 14M (Table 8). It is not the fixed- P_h comparison shown in Figure 5.

Table 9: **31M results.** Base uses native PyTorch; T_2/P_h and T_7/P_h use the specialized implementation, with OL1 only at h and $\lambda = b = 1$. Each condition uses one training seed and full validation (Appendix C.3). Deltas use fixed Base loss 4.565514 and native latency 1.023861 ms; the same Base checkpoint takes 1.106255 ms under the kernel. $\mathcal{S}_{\text{model}}$ is measured by pooling integer logical-product counts, not an analytic ceiling. Timing follows Appendix E.2.

Recipe	κ	$\mathcal{S}_{\text{model}}$ (%)	Loss	Δ loss	Latency (ms)	Δ latency (ms)
T_0/P_0 (Base)	—	0.00	4.566	+0.000	1.024	+0.000
T_2/P_h	0	8.59	4.688	+0.123	1.021	-0.003
T_2/P_h	0.01	8.73	4.701	+0.136	1.017	-0.007
T_2/P_h	0.05	9.06	4.703	+0.137	0.969	-0.054
T_2/P_h	0.1	9.26	4.757	+0.192	0.908	-0.115
T_2/P_h	0.5	9.47	5.220	+0.655	0.681	-0.343
T_7/P_h	0	15.34	4.894	+0.329	0.987	-0.037
T_7/P_h	0.01	15.74	4.874	+0.309	0.985	-0.039
T_7/P_h	0.05	16.77	4.894	+0.328	0.896	-0.128
T_7/P_h	0.1	17.53	4.943	+0.377	0.801	-0.223
T_7/P_h	0.5	22.66	5.521	+0.955	0.607	-0.417

D.2 410M RESULTS

Table 10 reports all twelve trained 410M conditions. The 410M base-model loss exceeds 70M’s 4.100. All sizes receive 712 optimizer steps and 1.493B input tokens, corresponding to 106.14, 48.96, 21.20 and 3.68 tokens per parameter at 14M, 31M, 70M and 410M. The 410M schedule also has a lower peak learning rate (3×10^{-4} rather than 10^{-3}). The training trajectories are retained in Figure 6. We therefore treat 410M as a limited stress test. The main runtime comparison uses 14M, 31M and 70M; the extensive mechanistic intervention analysis remains at 14M.

Table 10: **410M fixed-token stress test.** All twelve trained conditions; pressure uses OL1 with $\lambda = b = 1$. Ceiling is $100\mathcal{S}_{\text{model}}^{\text{max}}$ at 2,048 tokens. This size has no P_h or unpressured T_4/T_7 runs.

Recipe	κ	Loss	Ceiling (%)	$\mathcal{S}_{\text{model}}$ (%)
T_0/P_0 (Base)	—	4.5475	0.00	0.011
T_1/P_0 (ReLU)	—	4.6513	24.93	21.050
T_4/P_{all}	0	5.6921		38.339
	0.01	5.6914		39.149
	0.05	5.6789	74.78	42.125
	0.1	5.6789		45.688
	0.5	5.1911		71.591
T_7/P_{all}	0	5.4263		41.121
	0.01	5.4331		41.478
	0.05	5.4679	87.25	44.053
	0.1	5.4908		48.998
	0.5	5.1209		80.616

E KERNEL IMPLEMENTATION AND VALIDATION

E.1 SPARSE EXECUTION MECHANISMS

The specialized implementations preserve trained weights and thresholding while exploiting zero activations. At h, z , empty activation tiles can be identified before loading the corresponding weights, avoiding both weight reads and matrix instructions; sufficiently sparse rows may instead use scalar arithmetic. In the paths summarized in Table 11, zero detection at a, m, q, k, v follows operand loading, so skipping avoids arithmetic but retains those reads.

Table 11: **Sparse execution rules in the 14M and initial 70M implementations.** MMA denotes a matrix multiply–accumulate instruction. Tile dimensions are tokens by input features. All zero checks execute during the forward pass.

Operation	Execution rule	Work avoided
QKV, FFN-up (a, m)	Check activation fragments after both operands are loaded.	Skip matrix instructions for empty fragments; retain operand reads.
FFN-down (h), attention-output (z)	Check 8×16 tiles before loading weights. Rows with at most two nonzeros use scalar arithmetic when numerical guards allow; other rows use the dense matrix path.	Empty tiles avoid weight reads and matrix instructions. Scalar rows read only weights paired with nonzero activations.
Attention (QK, PV)	Check operand fragments after loading them.	Skip matrix instructions if either fragment is zero; retain operand reads, masking and softmax.

In Table 4, *bypass* is omitted MMA instructions divided by issued plus omitted instructions, pooled across layers and all 338 validation sequences. It includes padded work and, at h, z , instructions replaced by scalar arithmetic. Bypass differs from $\mathcal{S}_{\text{model}}$, the model-wide fraction of zero-operand scalar products, and is not the fraction of latency saved.

The initial 70M port of the 14M kernel takes 3.325 ms on Base versus 1.611 ms natively. A subsequent optimization pass improves sparse h, z execution and dense components, reducing Base latency to 2.732 ms while leaving a 1.121 ms gap. The main figure retains this partially optimized measurement series; it does not substitute the separate-session controls below.

The 31M implementation ports that optimized 70M policy (Run045, `opt073`) to hidden width 256, FFN width 1,024 and head dimension 32. It uses paired LayerNorm, native dense input projections, h, z output projections with empty-tile skipping, scalar execution for rows containing at most eight nonzeros and a matrix-instruction fallback, dense FlashAttention, and a dense full-vocabulary output head. Dimension-specific tiling and masks are adapted without a new optimization search. Base and T_2/P_h use `opt073-31m-v1` (Run054); the gate-aware `opt073-31m-t7-v1` (Run055) retains the same policy while applying the seven-site thresholds, including a, m gates in the paired normalization. Table 11 describes the historical implementations, not every path in these later variants.

E.2 NUMERICAL VALIDATION AND TIMING PROTOCOL

Measurements use an RTX 5090, BF16, batch one, 2,048-token causal sequences, and full vocabulary logits for every token (PyTorch 2.11.0, Transformers 5.12.1, CUDA 12.8). This is full-sequence inference with fixed weights and shapes; cached decoding and concurrent serving are outside the evaluated scope. Final checkpoints were checked on all 338 validation sequences in three independent processes. Relative to the unmodified eager reference, acceptance required finite outputs, $|y - \hat{y}| \leq 0.25 + 0.02|\hat{y}|$ for every logit, relative ℓ_2 error at most 0.02 per sequence, and absolute mean validation-loss difference at most 0.001. Exact agreement in the 14M control below is stronger than these general tolerances. At 31M, all eleven final checkpoints pass these criteria on all 338 sequences in three processes each: 18 Base/ T_2/P_h processes in Run054 and 15 T_7/P_h processes in Run055 (33 total). This qualification applies to the measured checkpoints, not every possible weight or activation pattern.

PyTorch/SDPA and specialized implementations use CUDA graphs with identical inputs. We report geometric means of synchronized host timings over 64 fixed validation sequences, seven randomized paired passes and three independent processes. Timing includes embeddings, transformer blocks, final normalization, vocabulary logits, thresholding and all input-dependent sparsity checks. Compilation, static weight preparation, graph capture and equal input staging are excluded. Activations are recomputed on every forward pass.

Cross-recipe speedups are $t_{\text{PyTorch}, T_0/P_0}/t_{\text{specialized}, r}$, using a fixed native Base reference per size. Table 8 uses 0.65544 ms at 14M and 1.61105 ms at 70M; Table 9 uses 1.02386 ms at 31M. The specialized 14M Base takes 0.65157 ms (0.652 ms rounded), a different denominator from native Base. The 31M Base/ T_2/P_h and T_7/P_h timings come from separate sessions on different RTX 5090 devices, with the Run054 native Base held fixed for both; small cross-session differences should not be overinterpreted. The controls below compare execution configurations of the same checkpoint, with explicit denominators.

E.3 CONTROLLED 14M EXECUTION ABLATIONS

Table 12 isolates sparse h, z execution at T_2/P_h , $\kappa = 0.1$, holding checkpoint, thresholding, fusion and other paths fixed. Disabling it removes tile skipping and scalar substitution, retaining the fused kernel’s dense path. All six layers have identical h, z values and zero masks; logits agree exactly on all 338 sequences in three processes. These full-model controls do not separate memory-read, inspection and arithmetic costs.

Table 12: **14M T_2/P_h execution control at $\kappa = 0.1$.** All configurations retain thresholding. Full-model geometric-mean latencies follow the timing protocol above; sparse execution includes skipping and scalar substitution.

Configuration	h sparse execution	z sparse execution	Latency (ms)
A: full execution	On	On	0.5588
B: disable h	Off	On	0.6485
C: disable z	On	Off	0.5630
D: disable both	Off	Off	0.6415

Joint savings ($D - A$) are $82.7 \mu\text{s}$, or 12.9% of D (12.8–13.0% across processes). Disabling h adds $89.7 \mu\text{s}$; enabling z saves $4.2 \mu\text{s}$ with h on but adds $7.0 \mu\text{s}$ with h off, precluding additive attribution. (The unmodified kernel takes 0.5614 ms, 12.5% below D ; its $2.6 \mu\text{s}$ offset from A limits interpretation of the smaller z effect.)

The same checkpoint takes 0.7018 ms under PyTorch/SDPA: A reduces latency by 20.4%, with $D - A$ accounting for 57.9% of the gap. This includes other implementation differences; the share does not apply to the Base-checkpoint comparison.

E.4 70M EXECUTION CONTROL

Table 13 holds the 70M checkpoint, thresholding and other components fixed. Disabling sparse h, z execution removes skipping and scalar substitution, retaining the same fused kernel’s dense path. All configurations pass the full-validation criteria above; bitwise agreement was not established.

Table 13: **70M T_7/P_h execution controls at $\kappa = 0.5$.** Full-model geometric means with weights, thresholding and other optimized components fixed. The session differs from the main-text latency sweep.

Execution configuration	Latency (ms)
Optimized implementation, sparse h, z execution enabled	1.128
Same implementation, sparse h, z execution disabled	2.704
Native PyTorch h, z replacement	1.365

Sparse h, z execution reduces latency by 58.3% relative to the 2.704 ms dense path. Against the faster 1.365 ms native h, z replacement, the reduction is 17.4%. This comparison includes fusion and layout effects, not only sparsity exploitation.