

ACTIVATION SPARSIFICATION IN TRANSFORMERS: THRESHOLD PLACEMENT, PRESSURE PLACEMENT, AND SPARSE EXECUTION

Anonymous authors

Paper under double-blind review

ABSTRACT

Activation sparsity can accelerate language-model inference when specialized kernels skip computation associated with zero activations. Existing methods obtain sparsity through different combinations of training-time pressure (P) and thresholding (T), but their relative and combined effects have not been isolated under a common setup. We study these interventions under matched pretraining conditions across Pythia 14M and 70M models. For each combination of P and T , we derive an architecture-dependent maximum achievable sparsity ($S_{\text{model}}^{\text{max}}$), measure model-wide attained activation sparsity (S_{model}) and the final inference latency under a specialized sparse kernel. Results points to four different findings. First, **aggressive combinations of interventions can extend the attainable sparsity range**. The higher-thresholding ($\kappa = 0.5$) broader-site pressure, T_7/P_{all} , reaches 27.5% and 40.6% model-wide sparsity at 14M and 70M, respectively 92% and 82% of the architectural ceilings, but at the validation loss cost of 0.62 and 1.12. Secondly, **sparsity and speedup increase together within recipes**. Although not all sparsity translates into speedups, they correlate. Exploiting it requires specialized kernel development. For example, the sparser 14M recipe T_7/P_{all} achieves 27.48% sparsity but is slower than T_4/P_{all} at 12.71% (0.473 versus 0.459 ms). Third, we **observe a phenomenon of sparsity spillover**: FFN-hidden (h) pressure increases sparsity at the unpressured attention-output site (z) by 10.7 percentage points relative to no pressure, and **the biggest latency gains come from bypassing them**. Also, pressuring h alone is less harmful to validation loss than global pressure. We hypothesize that some sites on the transformer block are more susceptible to sparsification than others. Our main takeaway is that training interventions and deployment kernels must be co-designed so that the induced sparsity is both high and structurally exploitable at inference time.

1 INTRODUCTION

Activation sparsity can improve language-model inference efficiency when specialized kernels skip computations associated with zero activations. Recent work has shown that sparsity can be exposed in pretrained models through post-hoc thresholding or induced during training, in both cases with measured inference speedups (Liu et al., 2024; Cetin et al., 2026). These results raise a practical design question: how to effectively sparsify a dense model in activations into an exploitable structure for optimized inference?

Existing approaches induce activation sparsity in different ways. TEAL applies post-hoc clipping to pretrained models, ProSparse uses ℓ_1 regularization on feed-forward (FFN) activations during training, and Q-Sparse applies top-k sparsification to both FFN and attention activations (Liu et al., 2024; Song et al., 2024; Wang et al., 2024). While each approach is effective in its own setting, it remains unclear how they compare under a common setup.

To address this gap, we pretrain Pythia-family models from scratch (Biderman et al., 2023) to study the effects of (i) activation ℓ_1 pressure (P) and (ii) *thresholding* (T). We focus on Pythia-14M and use Pythia-70M as a scale check. While pressure adds a gradient component to induce smaller activation magnitudes, thresholding maps small values into exact zeros. Each one of them can be placed




Figure 1: **Quality, sparsity and execution on Pythia-14M.** (a) Final validation loss and (b) full-model latency versus S_{model} for the same 40 checkpoints: the Base Model (T_0/P_0), GeLU $\rightarrow$ ReLU (T_1/P_0), T_1/P_h , T_2/P_h with thresholds at h , z , and T_4/T_7 with no pressure (P_0 , dash-dot), pressure at h (P_h , dashed), or pressure at all thresholded sites (P_{all} , solid). Curves connect separate trained settings: $\lambda = .05, .1, .5, 1$ for T_1/P_h ; $\kappa = 0, .01, .05, .1$ for T_2/P_h ; and $\kappa = 0, .01, .05, .1, .5$ for T_4/T_7 . Dotted curves in (a) show post-hoc clipping of the Base/ReLU controls at $p = 0, .1, \dots, .9$; high-loss tails extend beyond the view. Panel (b) shows trained checkpoints only. Vertical dotted lines mark analytic sparsity ceilings. Latencies use RTX5090, batch one and 2,048-token sequences.

in different activation sites, giving a combinatorial design set of sparse interventions. For each combination of P and T , we derive an architecture-dependent maximum achievable sparsity ($S_{\text{model}}^{\text{max}}$), and measure model-wide attained activation sparsity (S_{model}) and the final inference latency under a specialized sparse kernel. Figure 1 presents sparsity vs. quality trade-offs and inference latency for Pythia-14M, across P/T combinations and baseline controls.

Across 40 14M-parameter and 22 70M-parameter training and evaluation runs, we identify several insights. First, **the T/P train-time interventions extend the attainable sparsity range relative to post-hoc clipping**, approaching the architectural ceiling for the more aggressive recipes, but at a cost in quality (Figure 3). Second, **sparsity and speedup increase together within recipes**, but a larger S_{model} does not directly translate into a larger speedup. The main requirement for lower latency is the location and structure of zeros and whether the kernel can exploit them, with most gains coming from skipped computation at h and z (Table 2). Third, **sparsity remains exploitable across the tested scales**: skipping computations can reduce latency at both 14M and 70M, but a sparse kernel does not transfer directly and requires adaptation to the larger model. Fourth, **local pressure at h provides a more favorable balance between quality, sparsity and latency than global pressure** when combined with broader thresholding. It lowers loss in all 16 matched comparisons at moderate thresholds, and latency in 18 of 20 across the full grid (Figure 5). Finally, **local pressure can induce global changes in sparsity**: at 14M and $\kappa = 0.05$, pressure at h increases sparsity at z , which receives no direct pressure (Figure 6). Taken together, these results indicate that activation sparsification should not be treated as an isolated training objective: training interventions and deployment kernels must be co-designed so that the induced sparsity is both high and structurally exploitable at inference time.

2 RELATED WORK

2.1 SPARSIFICATION INTERVENTIONS

Post-hoc thresholding removes small activations at inference-time of a pre-trained model. While CATS thresholds small-magnitude activations within FFNs (Lee et al., 2024), TEAL extends it to attention (Liu et al., 2024). These studies motivate training-time interventions to adapt activation distributions to sparse computation and improve the trade-off between quality and sparsity.

ReLU Strikes Back studies the effect of activations functions in language models (Mirzadeh et al., 2023), showing clear sparsity gains on the use of ReLUs. ProSparse combines progressive ℓ_1 pressure and thresholding at FFNs during continued training (Song et al., 2024). It does not, however, extend its interventions to attention and restrict the analysis to local FFN sparsity. Cetin et al. (2026) extends ProSparse setup to pre-training and propose specialized kernels. These studies show that training can reshape the quality–sparsity trade-off, but they’re restricted to a single local intervention at FFN sites.

Q-Sparse extend sparsification beyond FFNs. It applies top- k selection to both attention and FFN projection inputs (Wang et al., 2024), but with no ℓ_1 pressure. Its top- k operation introduces, however, additional overhead due to the sorting operation. Spark Transformer tries to address it with statistical top- k (You et al., 2025), but we take a simpler approach and use a constant thresholding non-linearity. We argue that when the intervention is introduced during training, the network can adapt itself. Additionally, we evaluate the complementary effect of regularization pressure on both attention and FFNs.

2.2 INFERENCE SPEEDUP

Turning activation sparsity into speedup requires a specialized kernel adapted to the model architecture. TEAL accelerates batch-one autoregressive decoding, reporting up to $1.8\times$ speedup at 50% sparsity (Liu et al., 2024). Cetin et al. (2026) address batched matrix multiplication with their TwELL kernel, and report $1.3\times$ speedup for batch inference. Spark likewise couples its sparsity mechanism with an implementation that reports decoding gains (You et al., 2025). In each setting, skipped computation and memory access must compesate the additional overhead introuced by the sparse operations. This development has been increasingly accelerated by the use of coding agents, through iterative kernel development with execution and profiling feedback (Ouyang et al., 2025; Luo et al., 2026). We take a related approach and use GPT-6 Astra do develop a specialized kernel based on our training interventions (OpenAI, 2026).

3 METHODOLOGY

3.1 SPARSIFICATION INTERVENTIONS

A sparsification intervention specifies where activations are thresholded (thresholding sites $\mathcal{N}$) and where activation pressure is applied (pressure sites $\mathcal{P}$). The sets $\mathcal{N}$ and $\mathcal{P}$ need not coincide, and the thresholding and pressure functions may take different forms. Together, these choices define a combinatorial space of training interventions. We restrict this space to keep the computational cost manageable while addressing the central research questions.

For thresholding, we avoid the overhead of exact top- k selection (Wang et al., 2024; You et al., 2025) and instead use one-sided and symmetric fixed-threshold operators $G_{+,\kappa}(x) = x\mathbf{1}\{x \geq \kappa\}$ and $G_{\pm,\kappa}(x) = x\mathbf{1}\{|x| \geq \kappa\}$, with $\kappa \geq 0$. The constant κ controls the strength of the intervention and is held fixed within each training condition. At $\kappa = 0$, $G_{+,\kappa}$ reduces to ReLU, whereas $G_{\pm,\kappa}$ reduces to the identity.

For activation pressure, ℓ_1 averages across targeted sites and layers the mean absolute activation at each site after any thresholding. Ordinary L1 pressure minimizes $\mathcal{L}_{\text{task}} + \lambda\ell_1$, where λ controls the pressure intensity (Song et al., 2024; Cetin et al., 2026). To limit conflict between the ℓ_1 pressure and the task update while reducing the need to tune an additional hyperparameter, we introduce orthogonal ℓ_1 pressure (OL1), inspired by gradient-conflict projection methods (Yu et al., 2020; Hsieh et al., 2024). OL1 adds a separately preconditioned pressure correction after the task-only AdamW update, projects away its component opposing the adaptive task direction, and caps its norm relative to that direction. At a fixed optimizer state, once the norm budget is active, further increasing λ leaves the pressure correction unchanged up to numerical stabilizers. To keep the multisite intervention study tractable, we use a common $\lambda = 1$ and budget $b = 1$ and directly measure whether the budget becomes active during training. Appendix A.1 gives the full formulation, a brief comparison and shows that budget saturation depends strongly on the target set.

Figure 2 locates the intervention sites within the Pythia architecture. The variables a and m denote the normalized inputs to the attention and FFN branches, respectively; h is the FFN hidden activa-

tion; z is the concatenated attention context before the output projection; and q , k , and v are the query, key, and value operands. Attention probabilities are denoted by p . We apply $G_{+, \kappa}$ at selected sites in $\{a, m, h, z\}$ and $G_{\pm, \kappa}$ at selected sites in $\{q, k, v\}$. When h is targeted, the thresholding nonlinearity replaces GELU.

Table 1 summarizes the studied thresholding and pressure assignments and their sparsity ceilings across model sizes. Section 3.2 defines these ceilings.

Table 1: Thresholding, pressure and sparsity ceilings. T_1 replaces GeLU with ReLU at h . Pressure uses OL1: $P_1 = P_h$ targets h , and P_{all} matches the thresholding sites ($\mathcal{P} = \mathcal{N}$). For T_4/T_7 , threshold strength is swept over $\kappa \in \{0, 0.01, 0.05, 0.1, 0.5\}$, with OL1 pressure weight $\lambda = 1$ and task-relative norm budget $b = 1$ fixed. Only the 14M T_1/P_1 ablation sweeps $\lambda \in \{0.05, 0.1, 0.5, 1\}$ at $b = 1$. Ceilings are $100 \cdot \mathcal{S}_{\text{model}}^{\max}$ at 2,048 tokens; executed recipes by model size are listed in Appendix C.3.

Recipe	Threshold sites $\mathcal{N}$	Pressure sites $\mathcal{P}$	Ceiling (%)		
			14M	70M	410M
T_0/P_0 (Base Model)	$\emptyset$	$\emptyset$	0.00	0.00	0.00
T_1/P_0	h	$\emptyset$	4.28	12.35	24.93
T_1/P_1	h	h	4.28	12.35	24.93
T_4/P_0	a, m, h, z	$\emptyset$	12.83	37.06	74.78
T_4/P_h	a, m, h, z	h	12.83	37.06	74.78
T_4/P_{all}	a, m, h, z	$\mathcal{N}$	12.83	37.06	74.78
T_7/P_0	a, m, h, z, q, k, v	$\emptyset$	29.95	49.42	87.25
T_7/P_h	a, m, h, z, q, k, v	h	29.95	49.42	87.25
T_7/P_{all}	a, m, h, z, q, k, v	$\mathcal{N}$	29.95	49.42	87.25




Figure 2: Intervention sites in a shared Pythia block. The attention and FFN branches use separate LayerNorms and join the residual path. Green nodes denote activation sites, and purple nodes denote learned weights. Table 1 specifies the thresholding and pressure assignments.

3.2 SPARSITY MEASURE AND ARCHITECTURAL REACH

Activation sparsity is the fraction of zero activations at a site. We define *model-wide sparsity* as $\mathcal{S}_{\text{model}} = N_0/N_{\text{model}}$, where N_0 counts scalar multiplications with a zero activation operand in Pythia-block matrix products, and N_{model} counts all scalar multiplications in those blocks and the final output projection. We aggregate counts over validation data, excluding masked attention pairs.

To quantify intervention potential a priori, we define the *model-wide sparsity ceiling* as $\mathcal{S}_{\text{model}}^{\max} = N_{\text{reach}}(\mathcal{N})/N_{\text{model}}$, where $N_{\text{reach}}(\mathcal{N})$ counts scalar multiplications guaranteed a zero activation operand if all selected thresholding sites are 100% sparse. Thus, $\mathcal{S}_{\text{model}}^{\max}$ is the maximum sparsity directly attributable to $\mathcal{N}$. It depends on architecture, selected sites, and sequence length. Appendix B gives the precise accounting and definitions.

4 EXPERIMENTAL STUDY

We pretrain Pythia-14M, 70M, and 410M from random initialization (Biderman et al., 2023) on MiniPile (Kaddour, 2023), using 1.493 billion input tokens per model and training for one full-pass. We evaluate validation loss and $\mathcal{S}_{\text{model}}$ over all validation sequences. The intervention recipes are listed in Table 1.

Experiments across model sizes. The 14M runs provide the most extensive controlled ablation, used to identify threshold and pressure regimes. The 70M runs test whether these regimes persist in a larger model. We also attempted to extend this analysis to 410M, but under the same token budget and number of optimization steps, this leaves the model substantially less trained. The baseline run with no intervention had a higher final loss and larger late-training gradient norm, which is evidence of a different training regime. We therefore treat the 410M result as a limited, partial scale-up, and focus the detailed intervention analysis on the 14M and 70M runs.

Within each model size and intervention, runs share initialization, data order and training budget. Appendix C.3 gives the full protocol. For comparison, we apply TEAL-style post-hoc magnitude thresholding (Liu et al., 2024) at a, m, h, z , with percentile targets $p = 0, 0.1, \dots, 0.9$ and site-layer thresholds calibrated separately on ten training blocks.




Figure 3: **Quality–sparsity trade-offs at 14M and 70M.** 36/22 trained conditions, with loss scales fitted separately to each size. Colors distinguish executed recipes; lines connect separate threshold settings (κ), or pressure weights (λ) for T_1/P_1 . Dotted paths apply post-hoc clipping at a, m, h, z to the fixed base and ReLU controls. Vertical guides mark the T_4 and T_7 analytic sparsity ceilings.

Train-time interventions extends the attainable sparsity range, close to the architectural ceiling for the more aggressive recipes, but at a quality cost. Figure 3 summarizes 58 trained conditions at 14M and 70M. At $\kappa = 0.5$, T_7/P_{all} reaches 27.48% and 40.60% model-wide sparsity: 91.8% and 82.2% of the respective ceilings. Validation loss increases by 0.62 and 1.12 relative to the same-size base model. At intermediate $\kappa \in 0.01, 0.05$, local pressure at h present a comparable sparsity at lower loss, compared to global pressure P_{all} . At Section 4.2 we do a comparative analysis on the effect of global vs. local pressure.

4.1 FROM SPARSITY TO SPEEDUPS

Agentic kernel development. To test whether learned sparsity reduces inference latency, we used Codex with GPT-6 Astra (OpenAI, 2026), to develop specialized GPU kernels under human guidance. We started from the kernels released by Cetin et al. (2026), although they required adaptation for Pythia architecture. The agent proposed, implemented and tested the kernels speed and accuracy. During the development phase, we sampled a subset from the training set and tested all kernel

changes on them, comparing the model’s output without the custom GPU implementation. We rejected kernels that provided a different model output response. Once we found a kernel that provided speedup for a single sparse model with the same response, we tested it for each training recipe. For each model, we measured inference time and output correctness over the validation set. Its speedup was recorded and included only if these checks passed. We then repeated this evaluation at 70M after adapting the kernels to its larger matrix dimensions. Although we are not specialized in GPU-kernel design, this workflow made direct execution tests practical. This case illustrates how coding agents can support systems development.

Sparsity and speedup increase together within recipes. For the implemented kernel, Figure 4 shows an approximately linear increase within each trained threshold/pressure family over the tested grid, although with different slopes across families. At 14M and $\kappa = 0.5$, for example, T_7/P_{all} is much sparser than T_4/P_{all} (27.48% versus 12.71%), yet slightly slower ($1.38\times$ versus $1.42\times$ base-model speedup). Although they correlate, a larger S_{model} does not directly translate into a larger speedup. The main requirement is the place and structure of zeros and wheater the kernel can exploit them. We investigate it on Section 4.3.

For exploitable structure, scale increases the opportunity. Figure 4 shows that the attainable full-model speedup increases with model size. Larger models combine larger matrix operations with higher sparsity ceilings, making sparse computation more effective at reducing end-to-end latency. Accordingly, the best observed speedup rises from $1.42\times$ at 14M to $2.08\times$ at 70M.




Figure 4: **From model-wide sparsity to measured speedup.** Pythia-14M (a) and 70M (b), with 22 trained checkpoints and 20 post-hoc clipping settings per size. Speedup divides the same-size base-model latency by each setting’s latency, using the final specialized implementation for both terms. Timings use RTX5090, BF16, batch one, 2,048 tokens and the full logit head. Sparsity uses full-validation pooled logical counts.

Post-hoc clipping does not yield comparable gains in this implementation. Every post-hoc clipping point remains below $1\times$ at 14M, and the first degradation point show the additional overhead cost. At 70M, aggressive clipping can exceed $1\times$, reaching $1.22\times$, but with substantial loss degradation. Here clipping uses a separate masking operations around the kernel, adding additional overhead because it has to iterate over activations on every forward pass. The resulting zeros still have to form groups that the kernel skipping rules can use. The overhead can outweigh the avoided products. However, we did not optimize a fused clipping implementation, so this result does not establish a general limitation of post-hoc sparsification, but **points into a requirement of pairing the sparse intervention with specialized implementation.**

4.2 PAIRED EFFECTS OF GLOBAL VS. LOCAL PRESSURE

One question raised during sparse intervention design is wheater **activation pressure should be applied over the entire network, or are there local sites more suceptible to sparsification?** To answer this, we compare thresholding accross multiple sites (T_4 and T_7) with global and local FFN hidden pressure (P_{all} and P_h) at Figure 5.




Figure 5: **Paired effects of broadening pressure from P_h to P_{all} .** Every point is P_{all} minus P_h at matched model size, threshold sites and κ , with $\lambda = b = 1$. Rows show 14M/70M; columns show loss change, sparsity change (pp), and full-model latency change (μs). Positive loss or latency differences favor P_h ; positive sparsity differences favor P_{all} .

Global pressure achieves higher sparsity at a greater quality cost, without translating this additional sparsity into lower latency. At moderate thresholds ($\kappa \leq 0.1$), local h-only pressure (P_h) provides a more favorable quality–sparsity trade-off than global pressure (P_{all}) in both model sizes: P_h lowers loss in all 16 matched comparisons at 14M and 70M (Figure 5), with reductions of 0.19–0.32 while sparsity differs by at most 3.36 percentage points. Although broader pressure can achieve higher sparsity, this gain does not translate into lower measured latency: h-only pressure is faster in 18 of the 20 matched pairs, including all ten 70M pairs. Thus, applying sparsity pressure globally trades model quality for additional sparsity without yielding a corresponding inference-time benefit.

Local pressure can produce nonlocal sparsity changes. Figure 6 present the sparsity by site and layer across T_4 and T_7 interventions, for all pressures at 14M. At $\kappa = 0.05$, P_h increases sparsity at both h and z relative to P_0 , although z receives no direct pressure. Under T_7 , P_h also produces higher sparsity at h and z than P_{all} , despite the latter targeting both sites. One possible explanation is that (q,k,v) increases ℓ_1 pressure norms, and P_{all} caps the OL1 norm more frequently (Appendix A.1), reducing the local at h . The takeaway is that a local site pressure may spillover and affect the distributions of other sites.




Figure 6: **Activation sparsity by site and layer in Pythia-14M.** The six threshold/pressure recipes at $\kappa = 0.05$ (top) and 0.5 (bottom). Each cell gives the percentage of zero activations over all complete validation sequences at one site and layer. Q/K are post-RoPE; all panels share a 0–100% color scale. Values are rounded to whole percentages, so 0 or 100 does not establish a completely dense or all-zero tensor.

(h, z) are more susceptible to sparsification. At higher threshold ($\kappa = 0.5$) (h,z) sites become fully sparse, even if no pressure is applied; while (q,k,v) sites are in the [14, 36] and (a,m) are in [57, 70] sparsity range. This implies the distribution of (q,k,v) and (a,m) may have different ranges and require different threshold values than (h,z).

4.3 WHEN SPARSITY TRANSLATES TO SPEEDUP

Sparsity reduces latency when avoided work exceeds the overhead cost of detecting and skipping it. At h, z, the sparse kernel checks 8×16 activation tiles before loading their associated weights, so empty tiles avoid both weight load from memory and the matrix multiplications. This strategy, however, doesn't extend to other sites. Applying it at a and m increased full-model latency by 17.2% and 27.4% on 14M T_7/P_{all} at $\kappa = 0.5$, with unchanged outputs. Only 11.12% of a's 8×16 tiles

were empty, and none at m , so most groups still required matrix multiplication and the additional overhead was not justified.

Attention follows FlashAttention’s tiled sequence (Dao et al., 2022): load q, k , compute scores, apply masking and softmax, then multiply probabilities by v . Our implementation add additional overhead after the checks, so skipping an empty tile’s multiplication cannot recover the latency from reading from memory. Softmax still runs and turns zero scores into nonzero probabilities, so q/k sparsity does not automatically propagate to pv . Zeros in v remain usable, and sparsifying the probabilities is a further untested opportunity.

We isolate each site’s effect at 14M by disabling it during inference time. Table 2 reports the *bypass* (i.e. percentage of skipped matrix instructions) and saved time per-site sparse operation for the sparser 14M recipe $T_7/P_{\text{all}}, \kappa = 0.5$. It makes clear that savings come from h (150 μs) and z (29 μs). Individual a, m, q, k effects remain unresolved against process variation, whereas pv skipping adds latency. Enabling qk/pv skipping together raises full-model latency from 0.514 to 0.521 ms despite substantial bypass. Thus, for this checkpoint and implementation, attention’s arithmetic savings do not offset its overhead, while h/z benefit from avoiding weight reads as well as computation.

Table 2: **Conditional latency savings:** 14M $T_7/P_{\text{all}}, \kappa = 0.5$. Latency is measured per 2,048-token sequence on RTX 5090 (BF16, batch one). Saved time is full-model latency with one skipping path off minus latency with all paths on; positive values indicate a benefit (reduced latency). Brackets give [minimum, maximum] savings across all off/on pairs of three independent timing runs. Each run reports the geometric mean over 64 sequences.

Operation	Bypass (%)	Saved time (μs)	Process span (μs)
a : QKV projection	8.4	−1.6	[−6.0, +3.8]
m : FFN-up	0.0	−2.4	[−5.1, +2.7]
h : FFN-down	94.5	+150.1	[+145.4, +154.4]
z : attention-output projection	99.6	+29.0	[+23.7, +35.7]
q, k : QK scores	58.0	−2.9	[−7.3, +1.1]
v : PV product	67.2	−6.2	[−9.0, −2.2]

5 DISCUSSION AND CONCLUSION

Our results suggest that sparsification should be understood as a joint training-and-deployment problem rather than as a property of activation distributions alone. First, **local pressure does not imply a local response**: applying pressure only at the FFN hidden site can alter sparsity at other sites, showing that the network can reorganize its internal representations in response to a localized training intervention. More broadly, **sparsity is useful only if the resulting zero structure can be exploited by the inference system**. S_{model} is best viewed as a training-time proxy for computational opportunity, not as a direct predictor of latency. In our measurements, the latency gains comes from (h, z) sites, which were the ones more susceptible to sparsification, with higher achieved sparsity. One hypothesis is that we must achieve a high ($> 90\%$) level of site-layer sparsity, so the zeros gets a exploitable tiled structure with higher probability.

Transformers appear sufficiently flexible to adapt their internal representations to architectural and optimization constraints, suggesting that larger gains may come from **designing training objectives around inference requirements**, rather than first maximizing a generic sparsity metric and only afterward attempting to accelerate it. The same principle may extend beyond activation sparsity: training could be used to shape representations for properties required by weight sparsity, quantization, batched execution, or other hardware-specific constraints. In this view, deployment requirements are not merely constraints imposed after training, but potential signals for deciding what properties the network should learn.

Several limitations define the scope of these conclusions and point to concrete directions for further work. Our scaling evidence is limited: although the 410M experiments demonstrate that high model-wide sparsity remains attainable, the fixed token budget leaves 410M in a substantially different optimization regime, so it should be interpreted as a stress test rather than a clean scaling experiment.

Our runtime study is also restricted to batch-one, full-sequence inference on an RTX5090; we do not study how sparsity interacts with batched inference, cached decoding, or other hardware.

In addition, the full intervention space is combinatorial, and under the available compute budget we evaluate only a selected set of threshold and pressure placements rather than an exhaustive factorial study. The nonlocal response visible in Figure 6 suggests one particularly promising direction: interventions concentrated at h and z may preserve much of the execution benefit while avoiding pressure or thresholding at sites whose sparsity is expensive in quality or unprofitable for the current kernel. Related extensions include learning or tuning different thresholds κ at different sites and directly sparsifying attention probabilities after softmax, which could expose work in pv without relying on q/k zeros to propagate through softmax. More generally, rather than asking where the model can be made sparse, a productive systems-guided approach is to decompose full-model inference into individual operations, identify where latency and removable work are concentrated, determine what structured representation would allow those operations to be skipped efficiently, and then design training-time interventions that encourage precisely that structure.

REFERENCES

- Stella Biderman, Hailey Schoelkopf, Quentin Gregory Anthony, Herbie Bradley, Kyle O’Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. Pythia: A Suite for Analyzing Large Language Models Across Training and Scaling. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 2397–2430. PMLR, 2023. URL <https://proceedings.mlr.press/v202/biderman23a.html>.
- Edoardo Cetin, Stefano Peluchetti, Emilio Castillo, Akira Naruse, Mana Murakami, and Llion Jones. Sparser, Faster, Lighter Transformer Language Models. *arXiv preprint arXiv:2603.23198*, 2026. URL <https://arxiv.org/abs/2603.23198>.
- Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. *arXiv preprint arXiv:2205.14135*, 2022. URL <https://arxiv.org/abs/2205.14135>.
- Yu-Guan Hsieh, James Thornton, Eugene Ndiaye, Michal Klein, Marco Cuturi, and Pierre Ablin. Careful with that Scalpel: Improving Gradient Surgery with an EMA. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 19085–19100. PMLR, 2024. URL <https://proceedings.mlr.press/v235/hsieh24a.html>.
- Jean Kaddour. The MiniPile Challenge for Data-Efficient Language Models. *arXiv preprint arXiv:2304.08442*, 2023. URL <https://arxiv.org/abs/2304.08442>.
- Donghyun Lee, Je-Yong Lee, Genghan Zhang, Mo Tiwari, and Azalia Mirhoseini. CATS: Contextually-Aware Thresholding for Sparsity in Large Language Models. In *Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=v3w2a7EInO>.
- James Liu, Pragaash Ponnusamy, Tianle Cai, Han Guo, Yoon Kim, and Ben Athiwaratkun. Training-Free Activation Sparsity in Large Language Models. *arXiv preprint arXiv:2408.14690*, 2024. URL <https://arxiv.org/abs/2408.14690>.
- Ilya Loshchilov and Frank Hutter. Decoupled Weight Decay Regularization. In *International Conference on Learning Representations*, 2019. URL <https://arxiv.org/abs/1711.05101>.
- Mao Luo, Hongbin Li, Feng Lin, Hanling Yi, and Zhe Huang. Agentic Kernel Optimization: Generating State-of-the-Art GPU Kernels Without Hand-Written CUDA. *arXiv preprint arXiv:2608.14560*, 2026. URL <https://arxiv.org/abs/2608.14560>.
- Iman Mirzadeh, Keivan Alizadeh, Sachin Mehta, Carlo C. Del Mundo, Oncel Tuzel, Golnoosh Samei, Mohammad Rastegari, and Mehrdad Farajtabar. ReLU Strikes Back: Exploiting Activation Sparsity in Large Language Models. *arXiv preprint arXiv:2310.04564*, 2023. URL <https://arxiv.org/abs/2310.04564>.

- OpenAI. GPT-6 Astra: A New Generation of Intelligence. OpenAI, September 2026. URL <https://openai.com/index/gpt-6-astra/>. Accessed September 19, 2026.
- Anne Ouyang, Simon Guo, Simran Arora, Alex L. Zhang, William Hu, Christopher Ré, and Azalia Mirhoseini. KernelBench: Can LLMs Write Efficient GPU Kernels? In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 47356–47415. PMLR, 2025. URL <https://proceedings.mlr.press/v267/ouyang25a.html>.
- Chenyang Song, Xu Han, Zhengyan Zhang, Shengding Hu, Xiyu Shi, Kuai Li, Chen Chen, Zhiyuan Liu, Guangli Li, Tao Yang, and Maosong Sun. ProSparse: Introducing and Enhancing Intrinsic Activation Sparsity within Large Language Models. *arXiv preprint arXiv:2402.13516*, 2024. URL <https://arxiv.org/abs/2402.13516>.
- Hongyu Wang, Shuming Ma, Ruiping Wang, and Furu Wei. Q-Sparse: All Large Language Models can be Fully Sparsely-Activated. *arXiv preprint arXiv:2407.10969*, 2024. URL <https://arxiv.org/abs/2407.10969>.
- Chong You, Kan Wu, Zhipeng Jia, Lin Chen, Srinadh Bhojanapalli, Jiaxian Guo, Utku Evci, Jan Wassenberg, Praneeth Netrapalli, Jeremiah J. Willcock, Suvinay Subramanian, Felix Chern, Alek Andreev, Shreya Pathak, Felix Yu, Prateek Jain, David E. Culler, Henry M. Levy, and Sanjiv Kumar. Spark Transformer: Reactivating Sparsity in FFN and Attention. *arXiv preprint arXiv:2506.06644*, 2025. URL <https://arxiv.org/abs/2506.06644>.
- Tianhe Yu, Saurabh Kumar, Abhishek Gupta, Sergey Levine, Karol Hausman, and Chelsea Finn. Gradient Surgery for Multi-Task Learning. In *Advances in Neural Information Processing Systems*, 2020. URL <https://arxiv.org/abs/2001.06782>.

A INTERVENTION DEFINITIONS

A.1 ACTIVATION PRESSURE AND OL1

Objective and update. For J activation tensors A_j at the sites and layers targeted by $\mathcal{P}$, each with n_j entries after any thresholding, define

$$\ell_1(\theta; \mathcal{P}) = \frac{1}{J} \sum_{j=1}^J \frac{1}{n_j} \sum_{i=1}^{n_j} |A_{j,i}|. \quad (1)$$

Each tensor has equal weight, so changing $\mathcal{P}$ changes both the targets and normalization. Ordinary L1 applies AdamW (Loshchilov & Hutter, 2019) to the globally clipped gradient of $\mathcal{L}_{\text{task}} + \lambda \ell_1$. OL1 separates task and pressure updates: task-only AdamW is followed by a pressure correction that removes the opposing component and limits its norm. Pressure therefore does not directly enter Adam’s moment estimates.

Both gradients are evaluated at the same parameters and averaged over the same microbatches. The clipped task gradient updates AdamW’s bias-corrected moments $\hat{m}_t, \hat{v}_t$; the pressure gradient $g_1 = \nabla_{\theta} \ell_1$ is unclipped and has no momentum. Using the same preconditioner for both directions gives

$$u = \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon_A}, \quad w = \frac{g_1}{\sqrt{\hat{v}_t} + \epsilon_A}. \quad (2)$$

Norms and inner products include parameters with both gradients and initialized Adam state. For $c = \langle u, w \rangle$, $q = \|u\|_2^2$ and stabilizer $\epsilon > 0$,

$$\tilde{w} = \begin{cases} w - \frac{c}{q + \epsilon} u, & c < 0 \text{ and } q > \epsilon, \\ w, & \text{otherwise.} \end{cases} \quad (3)$$

Aligned pressure is retained. For weight $\lambda \geq 0$ and relative norm budget $b > 0$,

$$r = \frac{\lambda \|\tilde{w}\|_2}{\|u\|_2 + \epsilon}, \quad s = \min \left\{ 1, \frac{b}{r + \epsilon} \right\} \quad (r > 0), \quad (4)$$

with $s = 1$ when $r = 0$. After task-only AdamW, each parameter group p receives $\theta_p \leftarrow \theta_p - \alpha_p \lambda s \tilde{w}_p$, where α_p is its learning rate. Before learning-rate scaling and ignoring stabilizers, the correction norm is $\min\{\lambda \|\tilde{w}\|_2, b \|u\|_2\}$. At a fixed optimizer state, increasing λ has no further effect once the cap is active.

OL1 combines conflict projection (Yu et al., 2020) and a smoothed task direction (Hsieh et al., 2024) in Adam’s preconditioned coordinates. The stabilizer makes orthogonality approximate. This calculation excludes group learning rates and decoupled weight decay; avoiding opposition to the Adam direction does not guarantee unchanged task or validation loss.

L1 versus OL1. At 14M T_1/P_1 , matched initialization, data order and training budget yield similar results (Table 3). Sparsity differs by less than 0.02 percentage points at every λ . OL1 has slightly lower loss at $\lambda \in \{0.05, 0.1, 0.5\}$ and higher loss at $\lambda = 1$. One seed cannot establish whether these small loss differences persist across seeds.

Table 3: **L1 and OL1 at T_1/P_1 on Pythia-14M.** Final validation loss and model-wide sparsity after 712 matched training steps; OL1 uses $b = 1$. Both metrics cover all 338 complete validation blocks (500 documents; 1,444 tail tokens excluded). One initialization per condition.

λ	L1		OL1	
	S_{model} (%)	Val. loss	S_{model} (%)	Val. loss
0.05	3.1416	5.2061	3.1452	5.1981
0.1	3.3363	5.1655	3.3386	5.1594
0.5	3.7876	5.1127	3.7680	5.1102
1	3.9493	5.1023	3.9384	5.1210

Update geometry and budget saturation. We measure $\rho_{\text{opp}} = \max(0, -c)/(q + \epsilon)$, the removed component’s norm relative to $\|u\|_2$ before capping. All recorded steps satisfy $q > \epsilon$; aligned steps contribute zero. Figure 7 shows ρ_{opp} and r/b for the 20 14M OL1 conditions in Figure 1, using all 712 updates per condition at $\lambda = b = 1$. The cap is active ($s < 1$) on 96.99% of T_7/P_{all} steps, compared with 0.51% for T_4/P_{all} and 0.25% for either P_h recipe. The diagnostics are consistent with stronger pressure–task conflict when attention sites are included, but do not isolate each site’s contribution.




Figure 7: **OL1 geometry depends on pressure placement at 14M.** (a) Median removed opposing component, $100\rho_{\text{opp}}$; bars span the 25th–75th percentiles over 712 training steps per condition, not confidence intervals. (b) Pre-cap norm ratio r/b : faint curves show individual thresholds, and bold curves their pointwise median. The dotted line marks the norm cap. Colors and line styles match Figure 1; both y-axes are logarithmic, and $\lambda = b = 1$ throughout.

B SPARSITY ACCOUNTING

B.1 ACTIVATION AND MODEL-WIDE SPARSITY

Activation sparsity is the fraction of zero entries at a site, pooling counts across layers and validation sequences before dividing.

Model-wide sparsity counts scalar multiplications with a zero activation input:

$$\mathcal{S}_{\text{model}} = \frac{N_0}{N_{\text{model}}}. \quad (5)$$

N_0 counts these multiplications within Pythia blocks; N_{model} counts all multiplications in those blocks and the final vocabulary projection. The latter contributes only to the denominator. Counts are pooled across layers and validation sequences; percentages are $100\mathcal{S}_{\text{model}}$.

For a block’s linear projection from d_{in} to d_{out} features over n rows, the denominator contribution is $nd_{\text{in}}d_{\text{out}}$. Each zero input contributes d_{out} to N_0 , one multiplication per output feature. Zero weights do not count as activation sparsity.

In attention, both operands are activations. We count each multiplication once if either or both are zero. For queries q , keys k , values v and attention probabilities p , one head in one sequence contributes

$$N_{0,\text{QK}} = \sum_{t=1}^T \sum_{u=1}^t \sum_{j=1}^{d_h} \mathbf{1}\{q_{tj} = 0 \vee k_{uj} = 0\}, \quad (6)$$

$$N_{0,\text{PV}} = \sum_{t=1}^T \sum_{u=1}^t \sum_{j=1}^{d_h} \mathbf{1}\{p_{tu} = 0 \vee v_{uj} = 0\}. \quad (7)$$

Here T is sequence length, d_h the head dimension, t the query position and u the key/value position. Restricting $u \leq t$ excludes masked future positions from both counts. We sum over heads, layers and validation sequences, using the actual operands after thresholding and intervening transformations.

Normalization, softmax, activation functions, additions and memory transfers are excluded, so $\mathcal{S}_{\text{model}}$ does not directly measure latency or speedup.

B.2 MODEL-WIDE SPARSITY CEILING

Assume all activations at the thresholding sites $\mathcal{N}$ are zero. If $N_{\text{reach}}(\mathcal{N})$ counts the multiplications that must then have a zero activation input, the ceiling is

$$\mathcal{S}_{\text{model}}^{\text{max}} = \frac{N_{\text{reach}}(\mathcal{N})}{N_{\text{model}}}. \quad (8)$$

The denominator is unchanged, and each multiplication is counted once. The numerator includes downstream activations forced to zero: setting $v = 0$ also makes $z = 0$, affecting both PV and the attention-output projection. Zero attention scores do not imply zero softmax probabilities.

The ceiling depends only on architecture, selected sites and sequence length. Natural zeros outside those sites can raise measured $\mathcal{S}_{\text{model}}$ above it; the ceiling does not predict the associated quality cost. Appendix C.2 gives the counts for each Pythia size.

C EXPERIMENTAL SETUP DETAILS

C.1 PYTHIA ARCHITECTURE AND INTERVENTION SITES

Table 4 gives the shapes and positions of the sites in Figure 2. Here T is sequence length, d hidden width, d_f FFN width and H the number of heads, each with dimension $d_h = d/H$. Sites and threshold are shared across layers.

Table 4: Activation shapes and intervention positions for one sequence.

Site	Shape	Following matrix product	Position
a	$T \times d$	QKV projection	After attention LayerNorm
m	$T \times d$	FFN up projection	After FFN LayerNorm
h	$T \times d_f$	FFN down projection	Activation-function output
q, k	$H \times T \times d_h$	Attention scores	After RoPE
v	$H \times T \times d_h$	Probability-value	After the QKV split
z	$T \times d$	Attention-output projection	Concatenated attention context

C.2 ARCHITECTURE-SPECIFIC SPARSITY CEILINGS

For one sequence, each Pythia block uses $3Td^2$ scalar multiplications in the QKV projection, Td^2 in the attention-output projection, Tdd_f in each FFN projection, and $dT(T+1)/2$ in each of QK and PV . With L layers and vocabulary size V_{vocab} , the totals are

$$\begin{aligned} C_{\text{block}} &= T(4d^2 + 2dd_f) + dT(T+1), \\ C_{\text{model}} &= LC_{\text{block}} + TdV_{\text{vocab}}. \end{aligned} \quad (9)$$

Sequence count cancels from the ceiling ratio. For $d_f = 4d$, per-layer counts in the ceiling numerator are 0 for T_0/P_0 , $4Td^2$ for T_1/P_0 and T_1/P_1 , $12Td^2$ for T_4 , and C_{block} for T_7 . Multiplying by L/C_{model} gives $\mathcal{S}_{\text{model}}^{\text{max}}$, independent of pressure placement.

Table 1 uses $V_{\text{vocab}} = 50,304$ and $(L, d) = (6, 128), (6, 512), (24, 1,024)$ for 14M, 70M and 410M, respectively, with the common sequence length in Table 5.¹

¹Pinned architecture configurations: 14M, 70M, and 410M.

C.3 TRAINING AND EVALUATION

All sizes use the Pythia-14M-deduped tokenizer and seeds of 1234 for initialization and data order, with one seed per condition. AdamW settings are shared: $\beta = (0.9, 0.95)$, $\epsilon = 10^{-8}$, weight decay 0.1 and gradient clipping at norm 1. Attention and hidden dropout are zero. Training uses 712 updates with 1% linear warmup followed by cosine learning-rate decay. Table 5 gives the learning rates, microbatch sizes and shared effective batch. Gradients accumulate over the listed microbatches before each optimizer update.

Table 5: **Training settings across model sizes.** Effective batch size equals microbatch size times gradient accumulation. The common training budget gives fewer tokens per parameter at larger sizes.

Setting	14M	70M	410M
Parameters	14,067,712	70,426,624	405,334,016
Peak learning rate	10^{-3}	10^{-3}	3×10^{-4}
Minimum learning rate	10^{-4}	10^{-4}	3×10^{-5}
Microbatch (sequences)	32	4	4
Gradient accumulation (microbatches)	32	256	256
Effective batch (sequences/update)	1,024	1,024	1,024
Sequence length	2,048	2,048	2,048
Tokens per parameter	106.14	21.20	3.68

All sizes include the T_0/P_0 and T_1/P_0 controls. The 14M study adds T_1/P_1 with L1 or OL1, and T_4 and T_7 with P_0 , P_h or P_{all} . At 70M, T_4 and T_7 use P_h or P_{all} ; at 410M, only P_{all} . Parameter grids are given in Table 1. We evaluate the final checkpoint, without validation-based selection. The main figures measure loss and model-wide sparsity on the same validation sequences (Appendix B).

Figure 8 shows the training trajectories. They do not establish convergence, and the lower learning rates at 410M limit comparisons that attribute loss differences to model size alone.




Figure 8: **Base-model (T_0/P_0) training across model sizes.** (a) Training loss and (b) gradient norm before clipping (logarithmic scale). Faint lines show individual updates; solid lines show centered nine-update means, with shorter windows at the ends. Gradient norms are computed after accumulation and are not normalized by parameter count.

D COMPLETE RESULTS AND SUPPORTING DIAGNOSTICS

D.1 DETAILED TRAINED RESULTS

Tables 6–8 report all 74 trained conditions, including settings outside the favorable quality–sparsity trade-off. Loss and model-wide sparsity use the same definitions and final checkpoints as the main figures, with one training seed and complete validation per condition. Values are rounded.

Table 6: **Additional 14M ablations.** Pressure at h (top) and thresholding without pressure (bottom), paired by λ or κ . Controls appear in Table 7.

λ	T_1/P_1 (L1)		T_1/P_1 (OL1)	
	Loss	$\mathcal{S}_{\text{model}}$ (%)	Loss	$\mathcal{S}_{\text{model}}$ (%)
0.05	5.2061	3.142	5.1981	3.145
0.1	5.1655	3.336	5.1594	3.339
0.5	5.1127	3.788	5.1102	3.768
1	5.1023	3.949	5.1210	3.938

κ	T_4/P_0		T_7/P_0	
	Loss	$\mathcal{S}_{\text{model}}$ (%)	Loss	$\mathcal{S}_{\text{model}}$ (%)
0	5.4705	7.212	5.4684	7.218
0.01	5.4665	7.414	5.4588	7.618
0.05	5.4341	8.206	5.4379	9.127
0.1	5.4196	8.953	5.4287	10.425
0.5	5.6597	10.216	5.7029	15.387

Table 7: **Matched 14M/70M results.** The same threshold/pressure setting at both sizes. Pressure uses OL1 with $\lambda = b = 1$.

Recipe	κ	14M		70M	
		Loss	$\mathcal{S}_{\text{model}}$ (%)	Loss	$\mathcal{S}_{\text{model}}$ (%)
T_0/P_0 (Base)	—	5.2086	0.000	4.0998	0.000
T_1/P_0 (ReLU)	—	5.2696	2.714	4.2227	10.066
T_4/P_h	0	5.2157	8.409	4.6065	25.792
	0.01	5.2045	8.586	4.6185	25.971
	0.05	5.1956	9.036	4.5964	26.566
	0.1	5.2287	9.343	4.7020	27.218
	0.5	5.7227	10.227	5.3057	29.236
T_4/P_{all}	0	5.4582	7.810	4.8053	25.673
	0.01	5.4583	8.587	4.8227	26.224
	0.05	5.4898	10.456	4.8739	28.490
	0.1	5.5483	11.394	4.9437	30.574
	0.5	6.0380	12.713	5.3895	35.596
T_7/P_h	0	5.2198	8.399	4.6508	25.786
	0.01	5.1981	8.856	4.6122	26.009
	0.05	5.1950	10.126	4.6670	26.895
	0.1	5.2374	10.904	4.7189	27.779
	0.5	5.7320	16.664	5.3372	32.233
T_7/P_{all}	0	5.4802	7.054	4.9412	23.562
	0.01	5.4758	7.723	4.9364	24.226
	0.05	5.4628	9.863	4.9636	26.524
	0.1	5.4295	11.797	4.9501	28.694
	0.5	5.8294	27.483	5.2160	40.602

D.2 POST-HOC CLIPPING AT 14M

Figure 9 extends the base/ReLU comparison in Figure 3 to the trained checkpoints with retained post-hoc sweeps. The detailed view resolves the trade-offs near the trained models; the full-range view shows the loss cost of aggressive clipping. These are changes to fixed models at evaluation, rather than additional training conditions. The clipping protocol is the same as in Section 4. No corresponding sweep was collected for T_4/P_h or T_7/P_h .




Figure 9: **Post-hoc clipping on trained Pythia-14M models.** Thirty trained checkpoints and ten clipping targets per checkpoint ($p = 0, 0.1, \dots, 0.9$). Both panels show the same settings, with different loss ranges. Prominent markers and recipe lines identify trained models; dotted paths with small open markers follow each fixed model through increasing clipping targets. Colors and recipe labels match the main figures, with the L1 comparison included explicitly. Each path starts at its measured $p = 0$ point. Lines connect evaluated settings, not fitted frontiers.

D.3 410M RESULTS

410M as a fixed-token stress test. Figure 10 retains all twelve trained 410M conditions and both complete control-clipping paths. At $\kappa = 0.5$, T_7/P_{all} reaches 80.62% model-wide sparsity, or 92.4% of the seven-site ceiling, at loss 5.121 versus the base model’s 4.547. This tests attainable sparsity under a larger architecture and the same token budget; it is not a quality-scaling experiment. The 410M seven-site ceiling is 87.25%, compared with 49.42% at 70M, so greater architectural reach also contributes to the larger observed model-wide sparsity.

The 410M base-model loss exceeds 70M’s 4.100. All sizes receive 712 optimizer steps and 1.493B input tokens, corresponding to 106.14, 21.20 and 3.68 tokens per parameter at 14M, 70M and 410M. The 410M schedule also has a lower peak learning rate (3×10^{-4} rather than 10^{-3}). These differences may contribute to the non-monotonic quality ordering, but their effects are not isolated. The training trajectories are retained in Figure 8. We therefore treat 410M as a limited stress test and use 14M/70M for detailed intervention and runtime analysis.

Table 8: **410M fixed-token stress test.** All twelve trained conditions; pressure uses OL1 with $\lambda = b = 1$. This size has no P_h or unpressured T_4/T_7 runs.

Recipe	κ	Loss	$\mathcal{S}_{\text{model}}$ (%)
T_0/P_0 (Base)	—	4.5475	0.011
T_1/P_0 (ReLU)	—	4.6513	21.050
T_4/P_{all}	0	5.6921	38.339
	0.01	5.6914	39.149
	0.05	5.6789	42.125
	0.1	5.6789	45.688
	0.5	5.1911	71.591
T_7/P_{all}	0	5.4263	41.121
	0.01	5.4331	41.478
	0.05	5.4679	44.053
	0.1	5.4908	48.998
	0.5	5.1209	80.616

Pythia-410M: fixed-token sparsity stress test




Figure 10: **410M quality-sparsity stress test.** All twelve trained conditions: Base, ReLU, and T_4/P_{all} and T_7/P_{all} at $\kappa = 0, .01, .05, .1, .5$. Dotted paths show all ten post-hoc targets for each fixed Base/ReLU control, including their full loss range. Vertical guides mark the T_4 and T_7 sparsity ceilings.

E KERNEL DEVELOPMENT AND EVALUATION

Implementation. The specialized forward pass preserves each trained model’s weights, thresholds, biases and residual connections. It combines CUDA, CUTLASS and FlashAttention components to fuse operations and skip matrix instructions whose activation operands are zero. It introduces no additional pruning. One kernel computes the two input LayerNorms from their shared input while retaining separate affine parameters and thresholding. QKV layout conversion, RoPE and the selected q, k, v gates are fused; q, k gates follow RoPE. The FFN-down and attention-output projections are evaluated jointly with their biases and residual addition. The final vocabulary projection remains dense. Table 9 specifies the sparse execution rules.

The scalar path requires finite weights and biases, with nonzero magnitudes between 2^{-50} and 2^{50} ; selected activations obey the same magnitude bounds. Other rows use the matrix path. The two sizes

Table 9: **Sparse execution rules.** Tile dimensions are tokens by input features. A matrix multiply-accumulate (MMA) instruction operates on $16 \times 8 \times 16$ elements (output rows, output columns, reduction dimension). Zero checks and any scalar replacement execute during every forward pass.

Operation	Execution rule	Work avoided
QKV, FFN-up (a, m)	Test each 16×16 activation fragment after both operands have been loaded.	Skip the MMA for an empty fragment; retain operand reads.
FFN-down, attention output (h, z)	Inspect eight rows; use scalar arithmetic for rows with at most two nonzeros when numerical guards pass. Process the remaining rows in 8×16 tiles padded to 16 rows.	Skip empty tiles before loading their weights. Scalar rows read only the corresponding weights.
Attention (QK, PV)	After loading the operands, use a warp-wide check to test whether either complete MMA fragment is zero.	Skip that MMA; retain input loading, masking, softmax and output processing.

retain these rules, with projection dimensions adapted from $(d, d_f) = (128, 512)$ to $(512, 2048)$. The 14M attention kernel uses 64×256 query/key tiles and two key/value partitions; 70M uses 128×128 tiles without partitioning to match its native attention reduction order. This change was needed for numerical agreement, rather than selected by a 70M performance search.

Scope and numerical checks. Tests use one RTX 5090, BF16 inference, batch size one, 2,048-token causal sequences, and all 50,304 output logits, with PyTorch 2.11.0, Transformers 5.12.1 and CUDA 12.8. The implementation assumes fixed weights, shapes and non-overlapping use of its buffers. It is not evaluated for training, cached decoding, concurrent serving, arbitrary masks or thresholds outside the tested grid. Each final checkpoint must pass in three fresh processes on all 338 validation sequences. Relative to the unmodified eager implementation, outputs must be finite, every logit must satisfy $|y - \hat{y}| \leq 0.25 + 0.02|\hat{y}|$, the relative ℓ_2 error of each sequence’s logits must not exceed 0.02, and the absolute mean validation-loss difference must not exceed 0.001. These are numerical tolerances, not bitwise equivalence.

Timing and reference. Native PyTorch/SDPA and specialized inference use CUDA graphs with identical inputs. Latency is the geometric mean of host timings over 64 fixed validation sequences, seven randomized paired passes and three fresh processes. Compilation, static weight layout preparation, graph capture and equal input staging are excluded. Embeddings, all blocks, final normalization, vocabulary logits and input-dependent preprocessing are included. Post-hoc masks also remain inside timing; no activations are cached. Measurements use separate GPU sessions, so common normalization does not remove session effects.

For comparisons across recipes, Figure 11 uses *one native base-model latency per size*: $t_{\text{native}, T_0/P_0} / t_{\text{specialized}, r}$ for recipe r . The native references are 0.65544 ms at 14M and 1.66388 ms at 70M. This differs from Figure 4, which uses the specialized base-model latencies, 0.65157 and 3.31325 ms. The best trained-recipe gains are therefore $1.427\times$ and $1.028\times$ over the native base, versus $1.418\times$ and $2.047\times$ over the specialized base. The latter 70M ratio reflects a slow dense path in the ported implementation and does not establish a twofold gain over native inference.

What MMA bypass measures. Bypass is the number of omitted MMA instructions divided by issued plus omitted instructions, pooling integer counts across layers and validation sequences. Projection totals pool the four projection operations; attention totals pool QK and PV . We keep the two groups separate because equal instruction counts need not imply equal costs. At h, z , bypass includes replacement by scalar arithmetic and counts padded rows; it is not a fraction of all arithmetic eliminated. Attention counts include causal padding; even the base model bypasses 10.42% of PV instructions at 14M and 5.15% at 70M. In contrast, $\mathcal{S}_{\text{model}}$ excludes masked pairs and includes the dense vocabulary projection in its denominator.

The figure’s projection scalar sparsity pools zero-activation products over all projection products in the actual BF16 operands. Both this measure and bypass cover all 338 validation sequences; timing uses the 64-sequence subset. Thus these diagnostics complement, rather than redefine, the canonical model-wide sparsity used in the quality figures.




Figure 11: **Scalar sparsity, instruction bypass and speedup over a fixed native base model.** Rows show 14M and 70M, each with 22 trained settings and 20 post-hoc control settings. Left: projection scalar sparsity versus projection MMA bypass. Middle/right: projection/attention bypass versus full-model speedup. Every speedup within a row uses the same native T_0/P_0 reference; the horizontal line is $1\times$. The Base model marker uses specialized inference, so it need not lie on that line. Lines connect measured thresholds, not fitted trends; dotted paths are post-hoc clipping. Counts are pooled, include the instruction-padding and scalar-substitution conventions above, and do not measure saved memory traffic.

Interpretation and controlled tests. Projection bypass tracks speedup in both sizes, but high scalar sparsity alone need not produce empty instruction tiles. Post-hoc controls show this gap and incur separate masking costs. Attention bypass is nearly constant across the trained grid except for T_7/P_{all} at $\kappa = 0.5$; this single exception cannot establish a general attention-speedup relation. At 70M it has more attention bypass but less native-base speedup than T_7/P_h ($0.952\times$ versus $1.028\times$). Bypass therefore describes exploitable structure without explaining all latency differences.

Table 2 provides the direct 14M test: disable one skipping path while keeping weights, gates and fusion fixed. Its min-max spans cover all pairs of three independent off/on process means; effects are conditional and non-additive. These tests do not separately time memory reads, zero checks or softmax. No corresponding operation-level ablation was run at 70M, so its scatter is descriptive evidence, not causal attribution.

Alternatives tested. Porting the h, z strategy to a, m was correct but slower in the tested 14M case (Section 4.3). The padded fallback issued $1.90\times$ and $2.00\times$ as many matrix instructions, respectively. Attention shortcuts using cumulative value sums for zero queries, disjoint query/key supports, or cheaper conditional execution either failed numerical checks or did not establish a reliable gain. Disabling attention skipping improved all 35 historical 14M comparisons. The retained implementation is therefore not latency-optimal; it preserves one common execution policy across recipes. Earlier detection that avoids attention operand reads, fused post-hoc masking and sparsification of softmax probabilities remain untested here.