




Pythia-14M | RTX 5090 | BF16 | batch 1 x 2,048 tokens | full logits

35 variants; 3 processes x 64 inputs x 7 paired passes; 35/35 new kernels qualified

(a) Full-model acceleration



(b) Skipping contribution (detail)



Bars: three-process min-max (not confidence intervals). Panel (b): expanded vertical scale.
Canonical R: source FP16 logical opportunity. Execution: BF16.