




Pythia-14M | K050 | RTX 5090 | BF16 | batch 1 x 2,048 tokens | full logits

35 variants; 3 processes x 64 inputs x 7 paired passes; 35/35 variants qualify

(a) Full-model acceleration



(b) Sparsity-aware paths (detail)



Bars: three-process min-max (not confidence intervals). Panel (b): expanded vertical scale.
FP16 logical R; BF16 execution. Both controls share norm fusion; h/z disabled-path MMA has row padding.