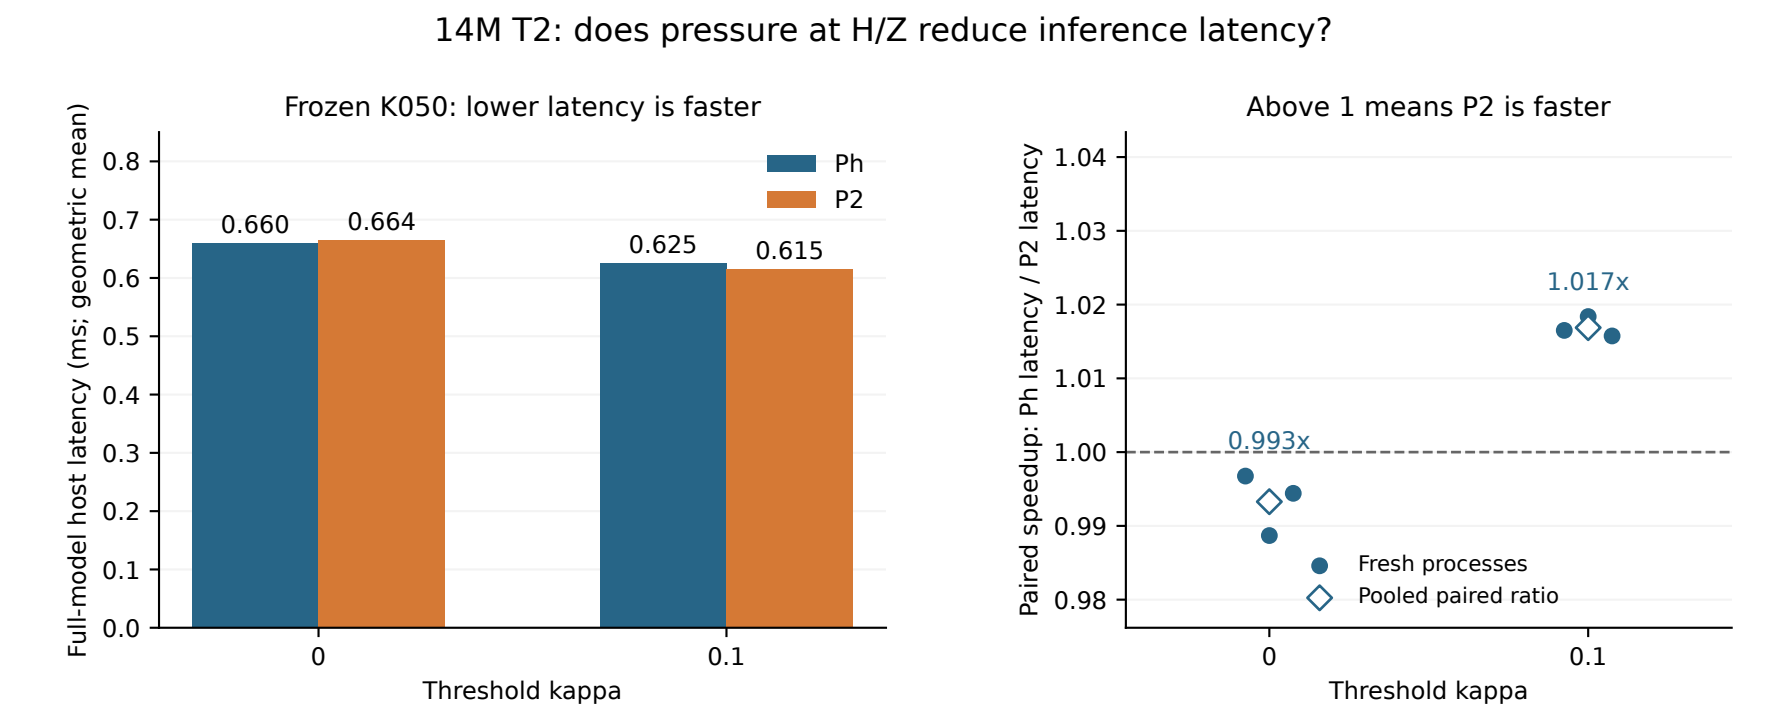
14M T2: does pressure at H/Z reduce inference latency?
Frozen K050: lower latency is faster
Full-model host latency (ms; geometric mean)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0.8
0.660
0.664
0.625
0.615
Ph
P2
0
0.1
Threshold kappa
Above 1 means P2 is faster
Paired speedup: Ph latency / P2 latency
0.98
0.99
1.00
1.01
1.02
1.03
1.04
1.017x
0.993x
Fresh processes
Pooled paired ratio
0
0.1
Threshold kappa
RTX 5090 | BF16 | batch 1 | 2,048-token uncached causal inference | all 50,304 logits
64 inputs x 7 paired passes x 3 fresh processes per threshold. All numerical checks passed.