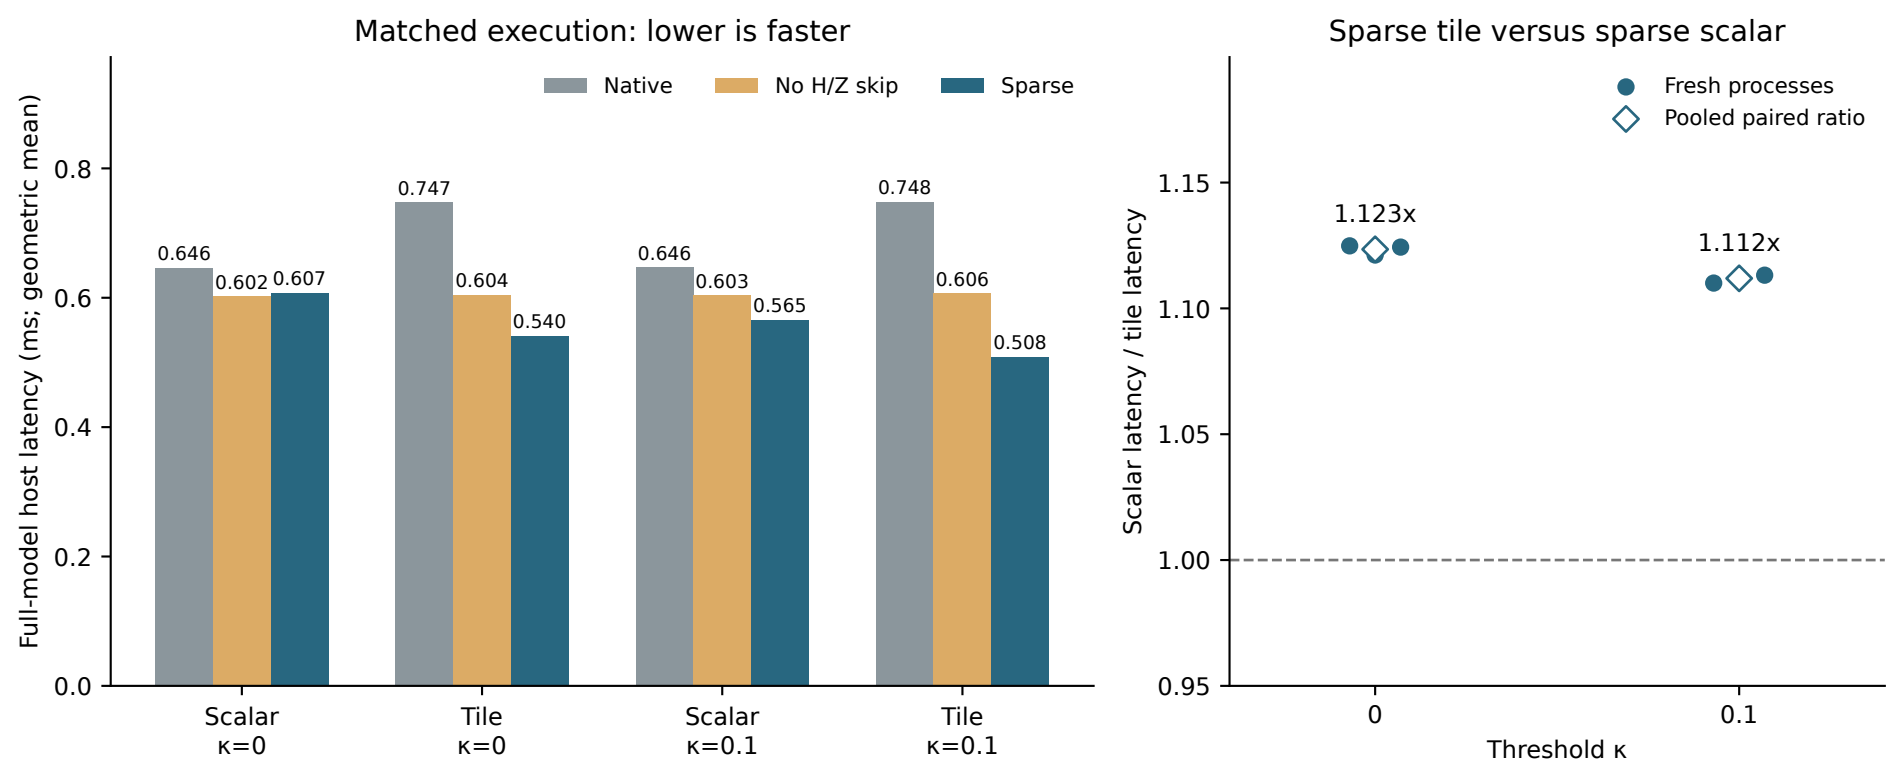
Matched execution: lower is faster
Native
No H/Z skip
Sparse
Full-model host latency (ms; geometric mean)
0.0
0.2
0.4
0.6
0.8
0.646
0.602
0.607
0.747
0.604
0.540
0.646
0.603
0.565
0.748
0.606
0.508
Scalar κ=0
Tile κ=0
Scalar κ=0.1
Tile κ=0.1
Sparse tile versus sparse scalar
Fresh processes
Pooled paired ratio
Scalar latency / tile latency
0.95
1.00
1.05
1.10
1.15
1.123x
1.112x
0
0.1
Threshold κ
RTX 5090 · BF16 · B1/T2048 · uncached · full logits · one training seed; different quality